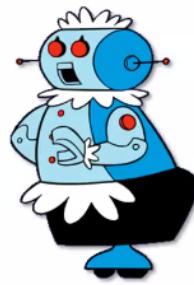
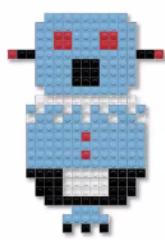


Aug-15

Speaker: Prof. Animesh Garg, University of Toronto

Title: Building Blocks of Generalizable Autonomy

Building blocks of Generalizable Autonomy in Robotics



Animesh Garg



Building blocks of Generalizable Autonomy in Robotics



Vacuuming



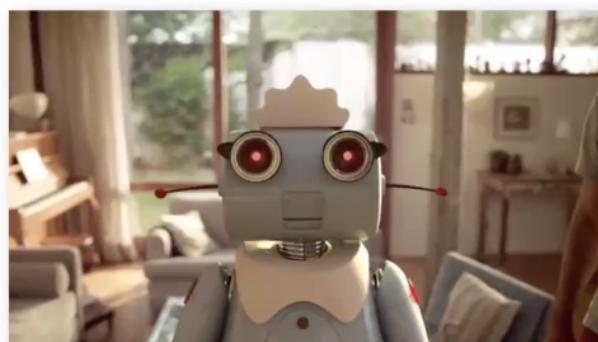
Sweeping/Mopping



Cooking



Laundry



Building blocks of Generalizable Autonomy in Robotics



Building blocks of Generalizable Autonomy in Robotics

How to achieve Algorithmic Generalization

Objects, Scenes, Semantics, Task Goals



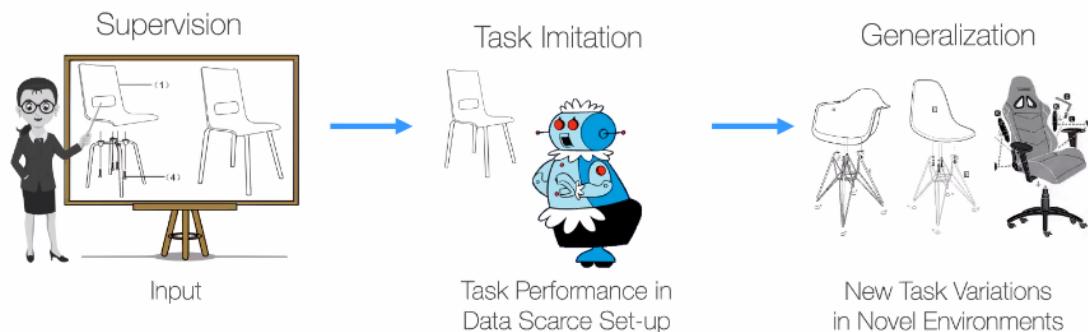
Known Structured Environment

Unstructured/Unknown New Environment

Boston Dynamics

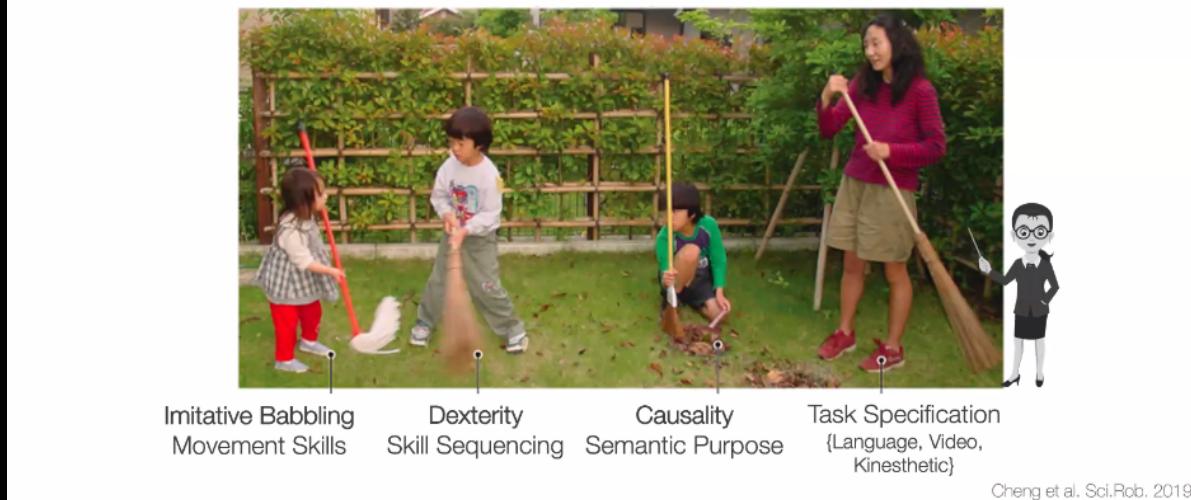
Building blocks of Generalizable Autonomy in Robotics

Vision: Build Intelligent Robotic Companions towards Human Enrichment and Augmentation

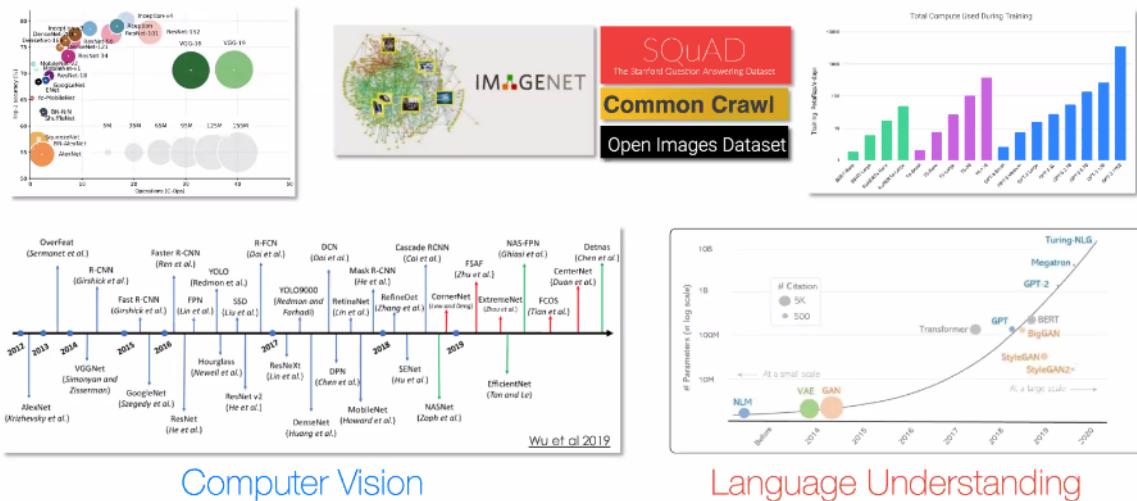


Building blocks of Generalizable Autonomy in Robotics

Imitation: But at which level? What should I copy?

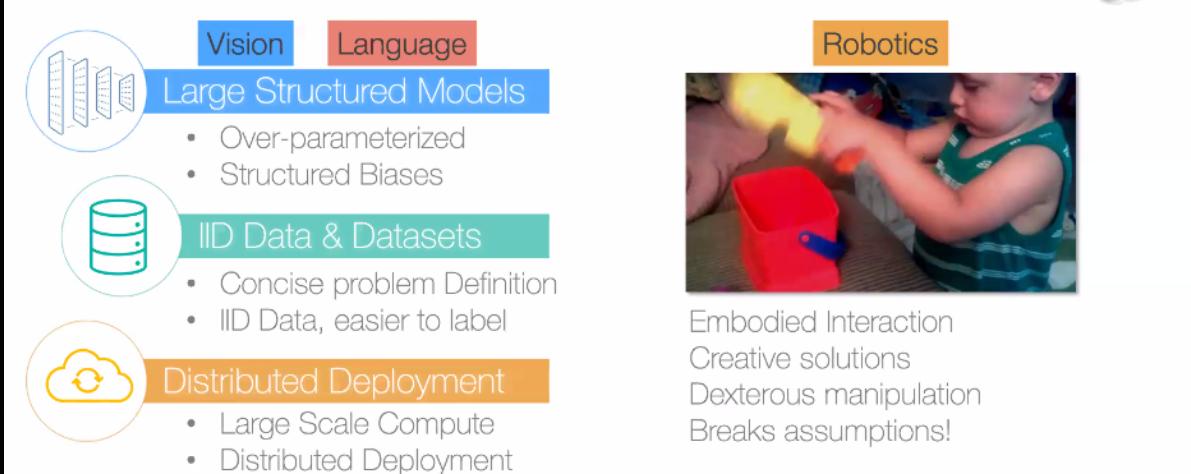


Generalization in Vision & Language



Generalization in Vision & Language

Ingredients of Modern Machine Learning & Applications



Generalization in Vision & Language

Ingredients of Modern Machine Learning & Applications

Physical Embodiment of Robots violates all three!



Vision Language Large Structured Models

- Over-parameterized
- Structured Biases

Robotics



IID Data & Datasets

- Concise problem Definition
- IID Data, easier to label

- X Small Model Sizes
X Lack of Structured Biases



Distributed Deployment

- Large Scale Compute
- Distributed Deployment

- X Variety of tasks
X Non-IID Sequential Data

- X Deployment is tied to training
X Low cloud compute usage

Generalizable Autonomy in Robotics

Vision: Build Intelligent Robotic Companions

Structured Representations and Causal Discovery



Structure

Representations for Robotics



Discovery

Data & Causality



Deployment

Transfer & Safety

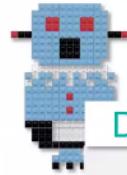


Building blocks of Generalizable Autonomy in Robotics



Structure

Representations for Robotics



Discovery

Data & Causality



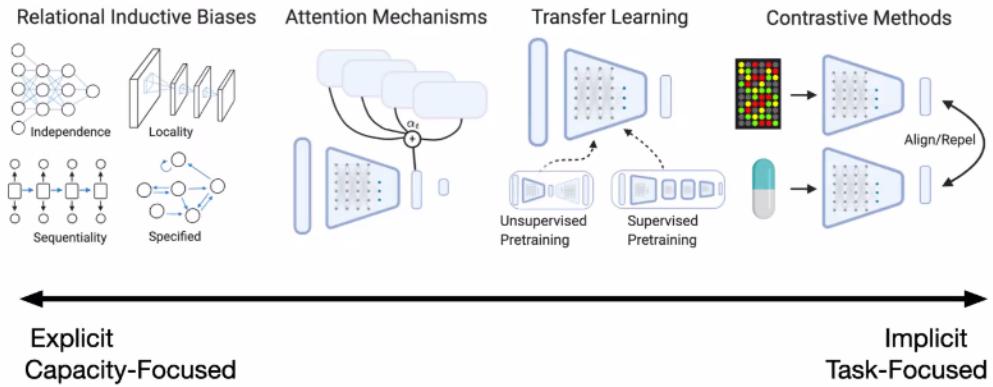
Deployment

Transfer & Safety

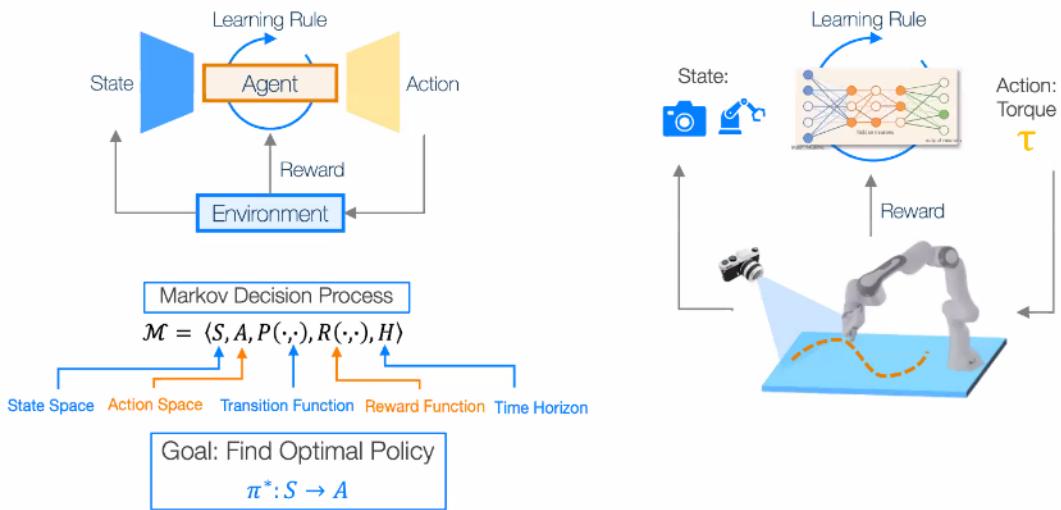


Representations in Machine Learning

Insight: Structure makes learning possible



Representations for Reinforcement Learning



Structure

State/Action Reps.
VICES IROS19
LASER ICR21
Making Sense ICRA19
Unsup KPs PAMI21

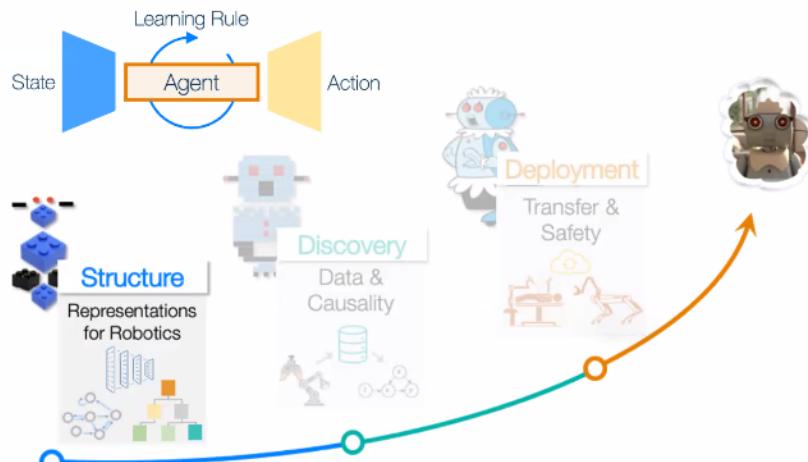
Inductive Biases
C-Learning ICLR21
OCEAN UAI20
D2RL arXiv20

Structure in Planning
CAVIN CORL20
Skill Hierarchy ICLR21
Finding-IT CVPR18

Neural Programming
NTP ICRA18
NTG CVPR19
Cont.Relax IROS19

Representations for Robot Learning

How can better state and action representations lead to generalization in Robotics?



Representations for RL Task Spaces

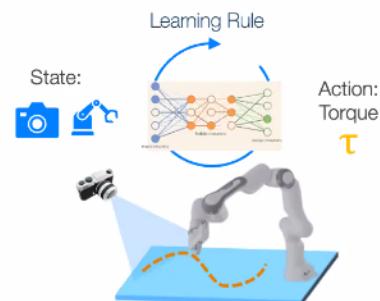
Latent (low-dim) manifold for maximum task-progress (reward)!



People Opening a Door for Friends
No vertical/sideways yanking (!)



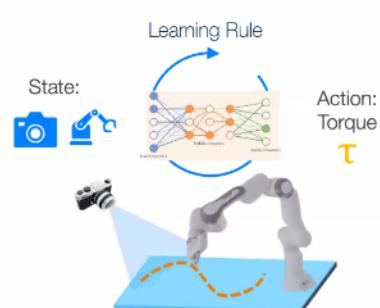
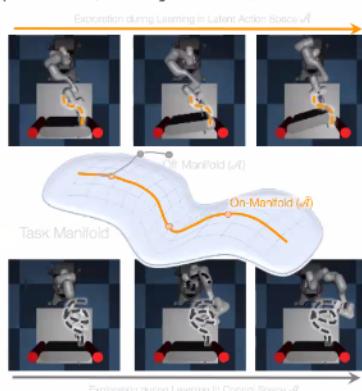
Robot Trying to Open a Door
Applies forces in all direction (!)



Representations for RL Task Spaces

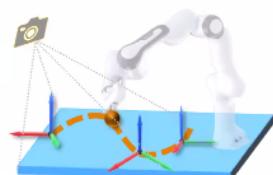
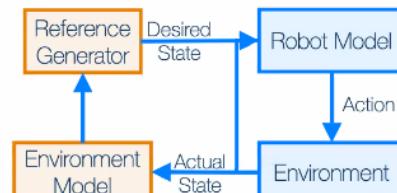
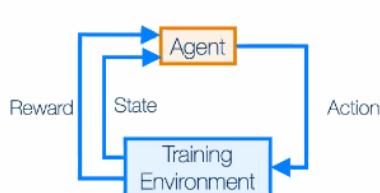
Latent (low-dim) manifold for maximum task-progress (reward)!

Multiple valid-task spaces, may exist, but only some are useful!



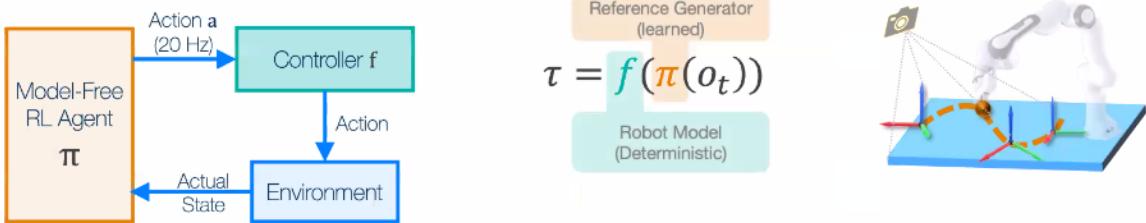
Representations RL: Task Spaces

RL with Variable Impedance Task-Space



Representations RL: Task Spaces

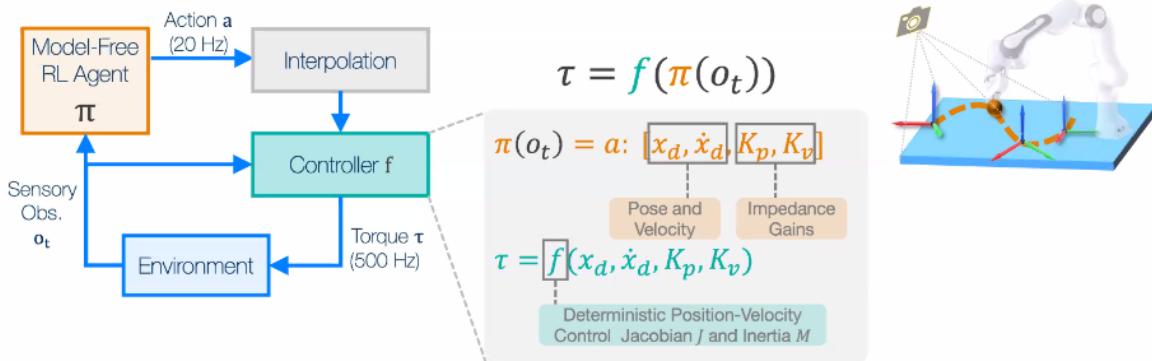
RL with Variable Impedance Task-Space



Variable Impedance Control in End-Effector Space, IROS 2019

Representations RL: Task Spaces

RL with Variable Impedance Task-Space



Variable Impedance Control in End-Effector Space, IROS 2019

Representations RL: Task Spaces

Surface Wiping

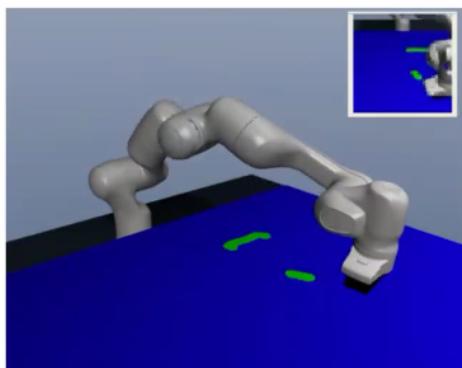
Input: Image (48x48)

Minimize the number
of Dirty Tiles

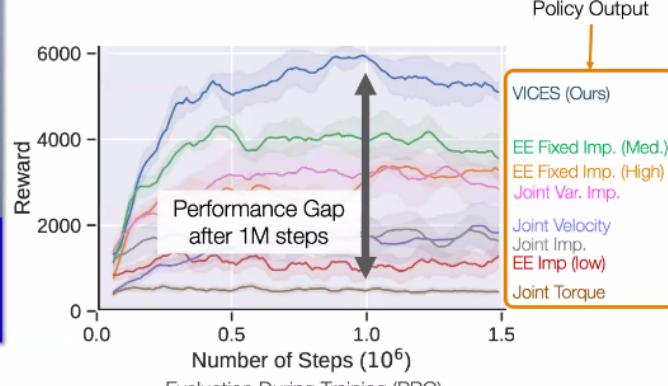
Reward: $\lambda_1 \sum(\text{dirt_on_table}) + \lambda_2(\text{distance_to_table}) - \lambda_3 \mathbb{I}(F \geq 40N)$

Maintain Contact
with the Table

Don't push with more
than Robot Payload



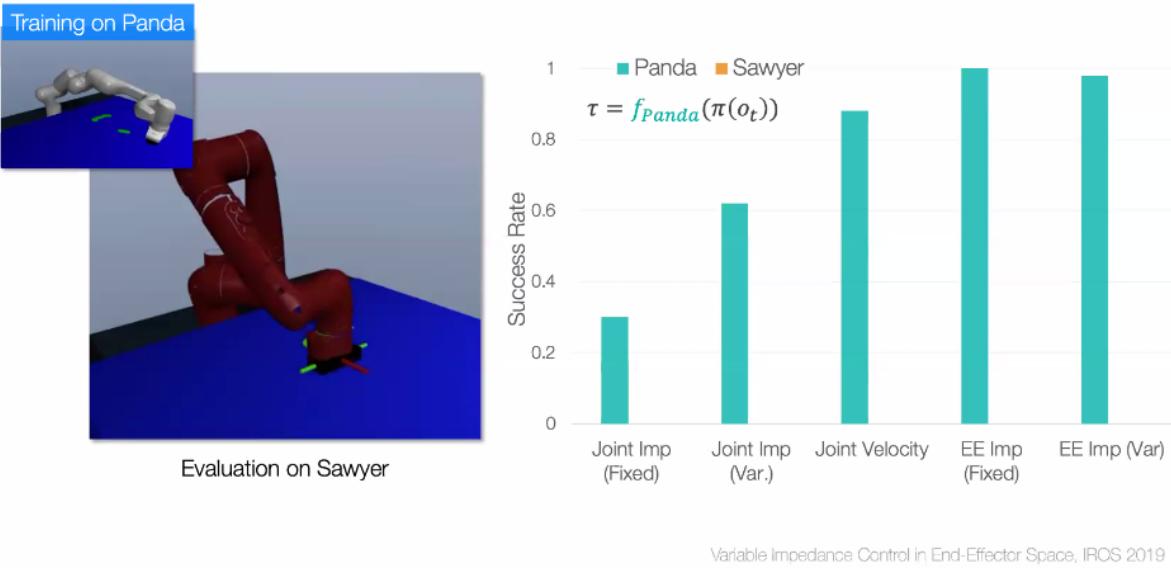
Trained Policy Rollout (Ours)



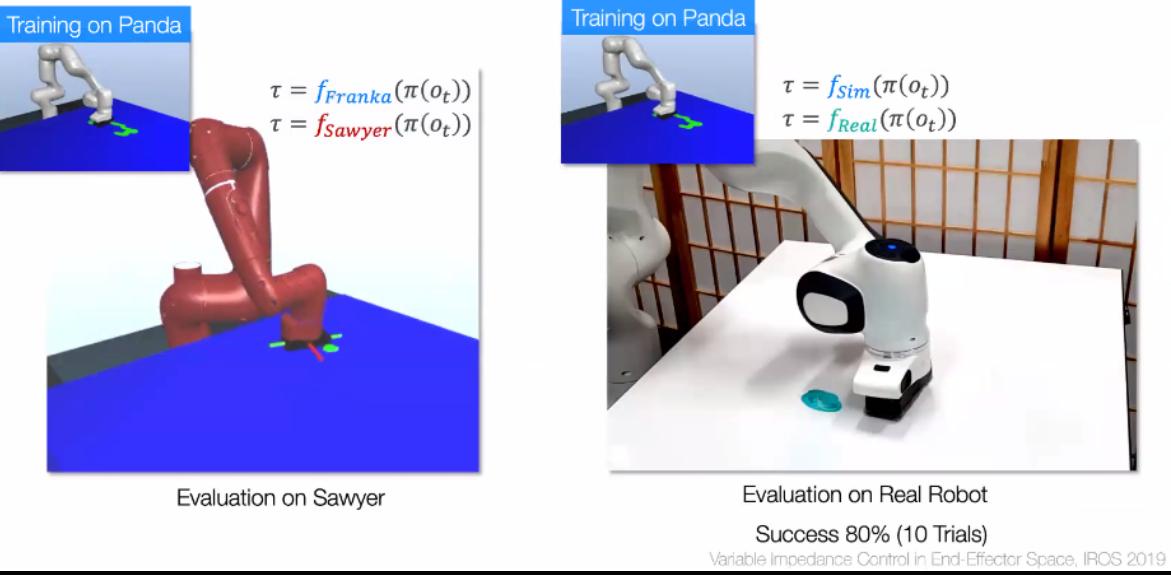
Evaluation During Training (PPO)

Variable Impedance Control in End-Effector Space, IROS 2019

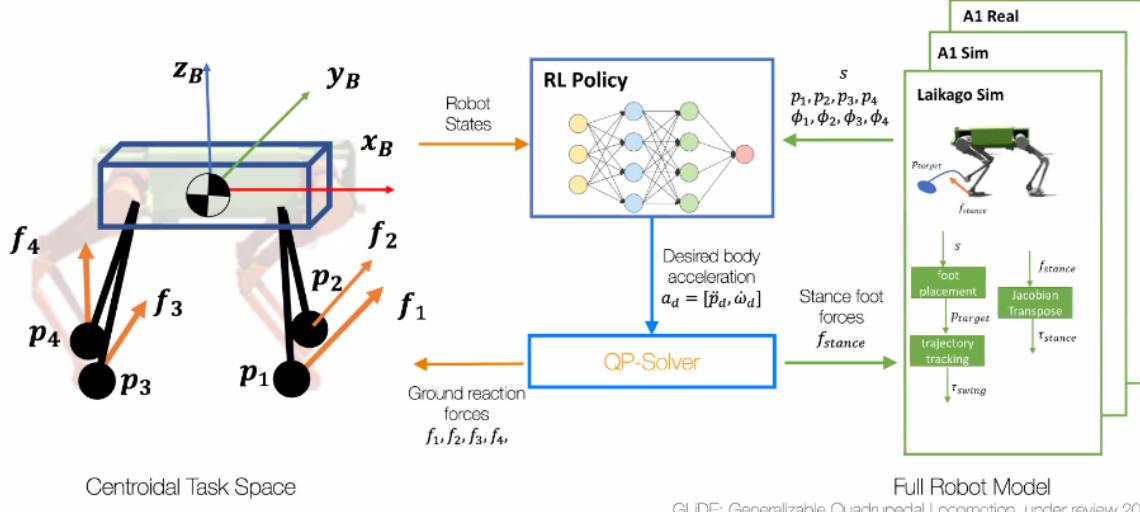
Representations RL: Task Spaces



Representations RL: Task Spaces

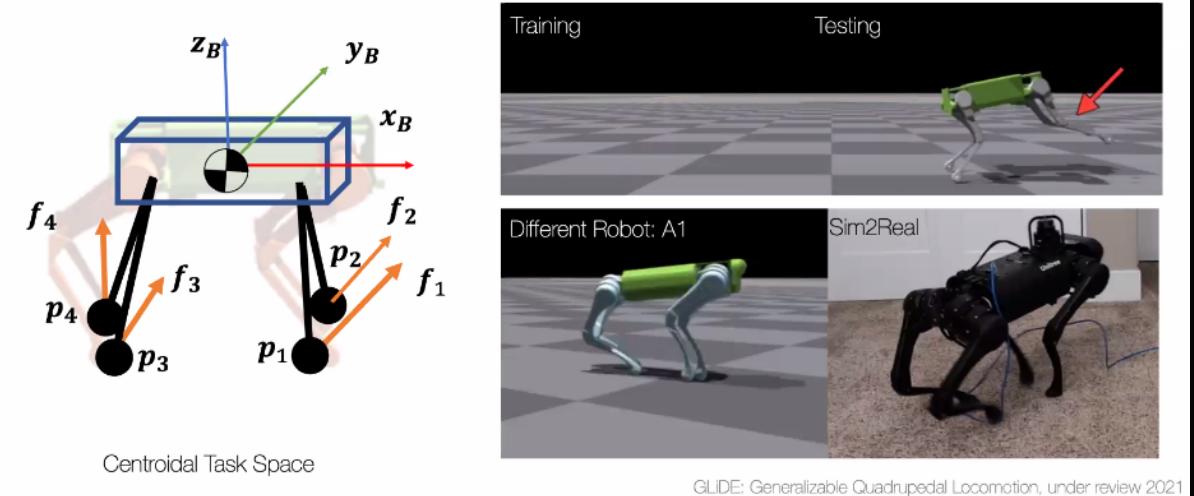


Representations RL: Task Spaces From manipulation to locomotion



Representations RL: Task Spaces

From manipulation to locomotion



Representations RL: LASER

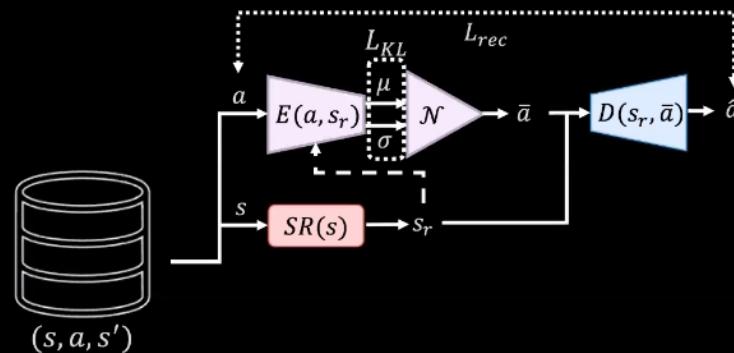
Learning the action space



LASER: Learning a Latent Action Space for Efficient RL, ICRA 2021

Representations RL: LASER

Learning the action space

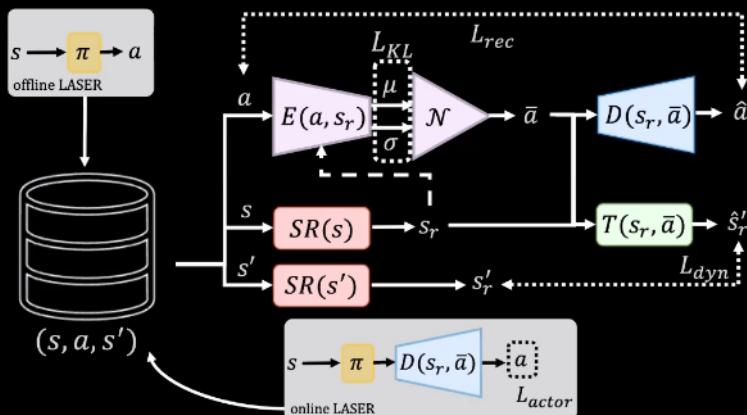


Variational auto encoder architecture with dynamic consistency

LASER: Learning a Latent Action Space for Efficient RL, ICRA 2021

Representations RL: LASER

Learning the action space



Can Use Batch Data as well as Replay Buffer concurrent with Policy

LASER: Learning a Latent Action Space for Efficient RL, ICRA 2021

LASER: Learning the action space Evaluation



SAC Baseline
(Same Task)



Offline LASER
(Same Task)

Train Action Space
then Train Policy

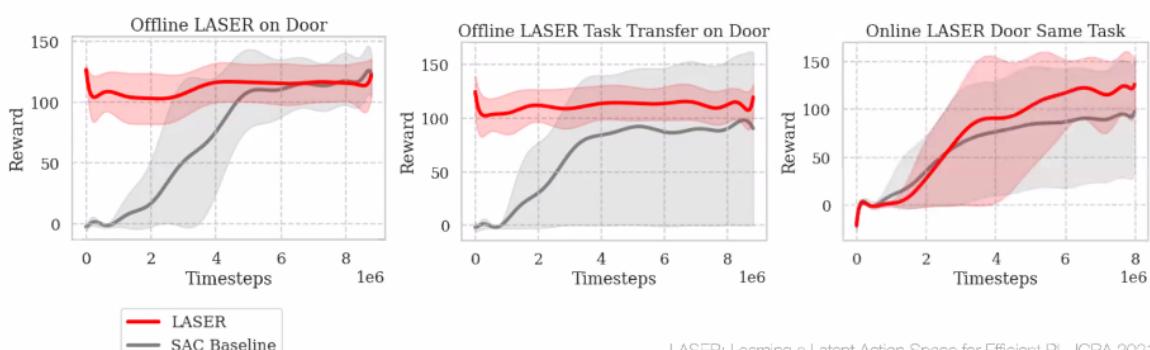
LASER: Learning a Latent Action Space for Efficient RL, ICRA 2021

LASER: Learning the action space Evaluation

Collect Task Demos,
Learn Action Space
[Does this work?](#)

It must have memorized the
solutions (data from experts)!
[Does it Generalize?](#)

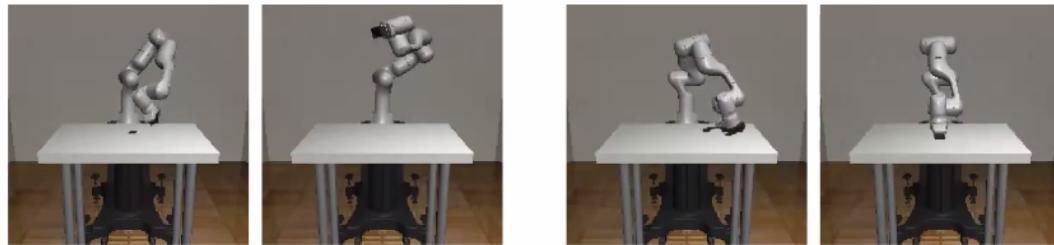
Gotcha! But policy already
known (good demos)!
[No expert\(s\), what then?](#)



LASER: Learning a Latent Action Space for Efficient RL, ICRA 2021

LASER: Learning the action space

Evaluation



SAC Baseline
(Same Task)

Offline LASER
(Same Task)

SAC Baseline
(Transfer Task)

Offline LASER
(Transfer Task)

Train Action Space
then Train Policy

Train Action Space
then Train Policy

LASER: Learning a Latent Action Space for Efficient RL, ICRA 2021

Representation: Action Spaces

Application: Assistive Teleoperation



Offline, personal demonstrations of high-dimensional motions



Learn low-dimensional latent representations for online control

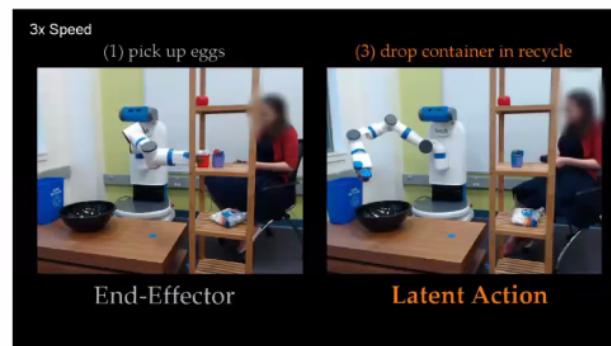


Easier to control **high-dimensional robots** by embedding the robot's actions into a **low-dimensional latent space**

Controlling Assistive Robots with Learned Latent Actions, ICRA 2020

Representation: Action Spaces

Application: Assistive Teleoperation

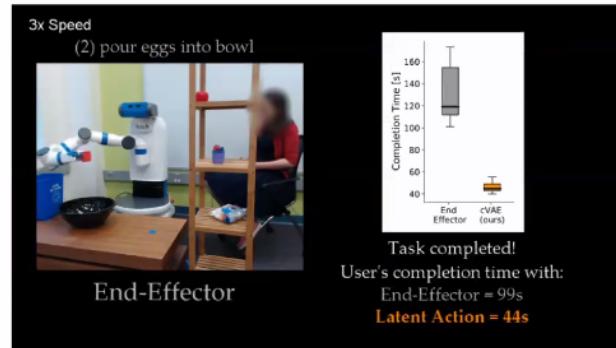
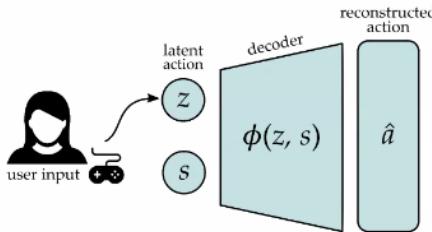


Easier to control **high-dimensional robots** by embedding the robot's actions into a **low-dimensional latent space**

Controlling Assistive Robots with Learned Latent Actions, ICRA 2020

Representation: Action Spaces

Application: Assistive Teleoperation



Easier to control **high-dimensional robots** by embedding the robot's actions into a **low-dimensional latent space**

Controlling Assistive Robots with Learned Latent Actions, ICRA 2020



Structure

State/Action Reps.
VICES IROS19
LASER ICRA21
Making Sense ICRA19
Unsup KPs PAMI21

Inductive Biases
C-Learning ICLR21
OCEAN UAI20
D2RL arXiv20

Structure in Planning
CAVIN CORL20
Skill Hierarchy ICLR21
Finding-IT, CVPR18

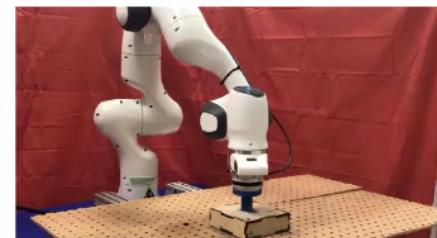
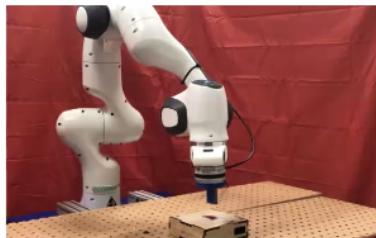
Neural Programming
NTP ICRA18
NTG CVPR19
Cont.Relax IROS19

Representations for Robot Learning

How can better state representations capture multimodal data?

Generalizable Multi-modal State Representations

- Learn a joint Visuo-Tactile representation for Peg Transfer
- Representation transfers to new task, while Policy doesn't



Making Sense of Vision and Touch, ICRA 2019 T2020

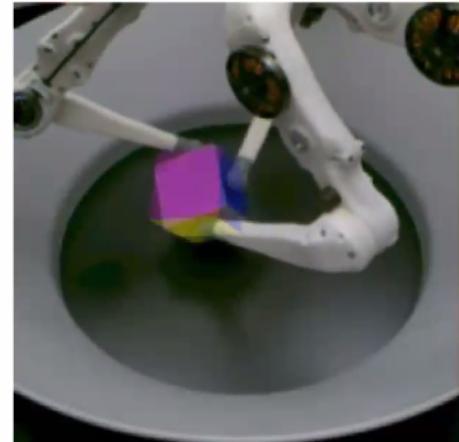
Representation: State Spaces

Application: Multi-finger Manip

Real Robot Challenge: Trifinger platforms

Task: repose in 6-DoF
(position + orientation)

Development done remotely in simulation
using Isaac Gym, no physical robot
access



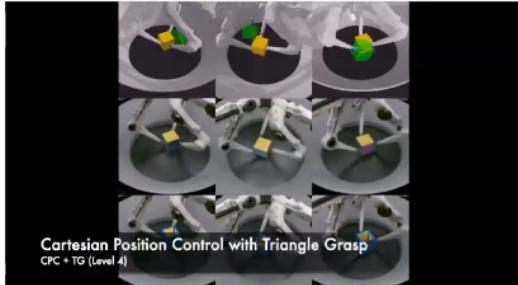
Easier to control **high-dimensional robots** with better action space and a **efficient physics sim**

CORL2021

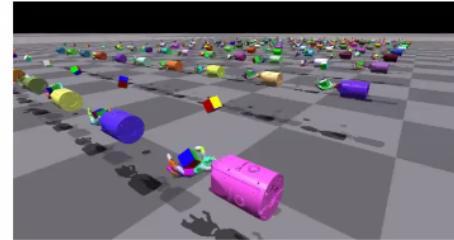
Background



Dexterous Manipulation via Simulation
OpenAI, 2018



Real Robot Challenge
Structured Policies

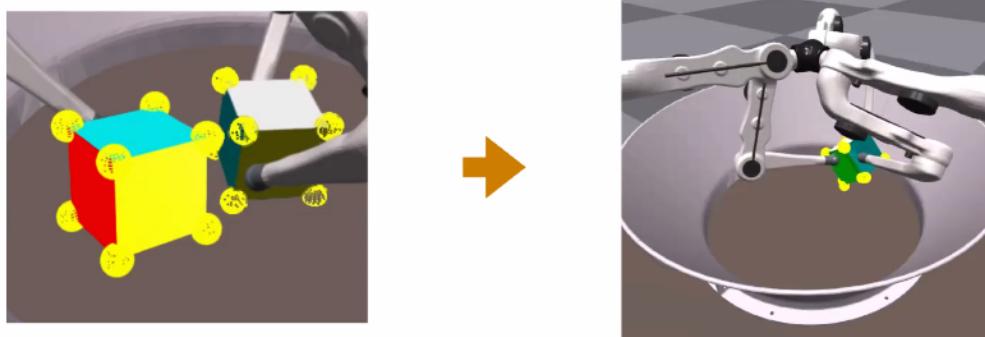


GPU-Simulated Manipulation
Isaac Gym

CORL2021

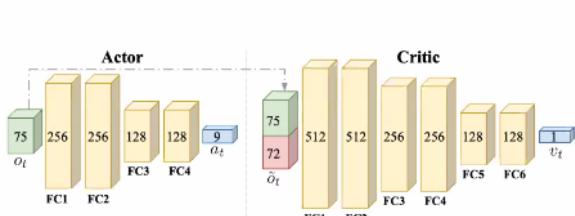
Reward and Pose Representation

- Traditional reward & observation performed poorly
- A better representation than position + quaternion in {observation, reward}?
- Allow for 6-DoF reposing



Representation: Action Spaces Application: Multi-finger Manip

Asymmetric Actor-Critic



- Policy Observations:
 - Joint Position
 - Joint Velocity
 - Object Pose
- Value function Observations:
 - Policy Observations
 - Joint Torques
 - Fingertip Positions
 - Fingertip Forces
 - Object Velocity



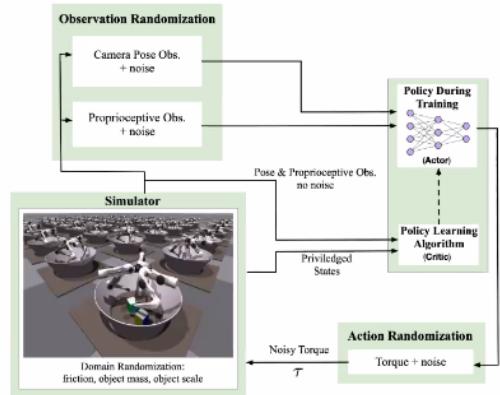
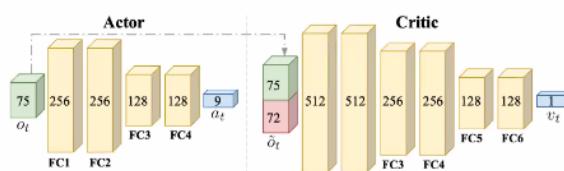
Easier to control **high-dimensional robots** by embedding the robot's actions into a **low-dimensional latent space**

CORL2021

Representation: Action Spaces

Application: Multi-finger Manip

Asymmetric Actor-Critic



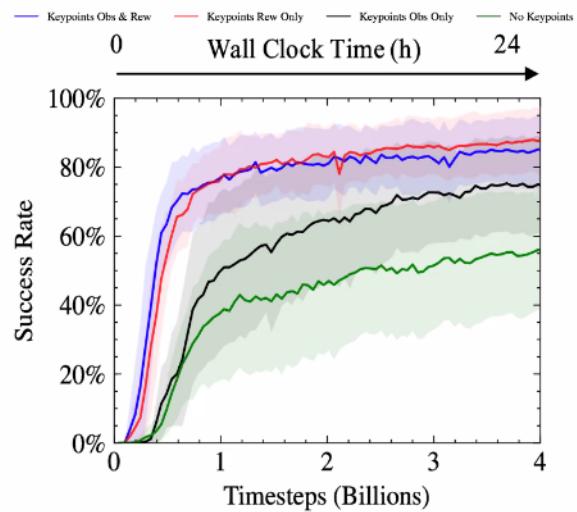
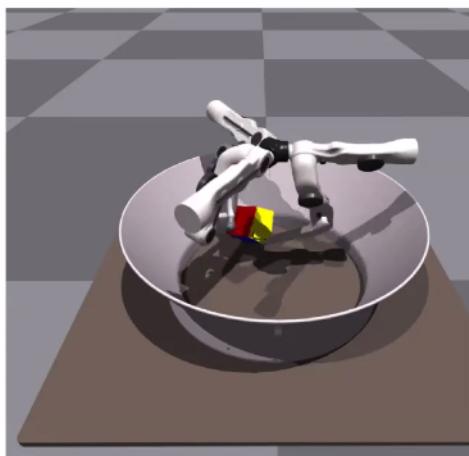
Easier to control **high-dimensional robots** by embedding the robot's actions into a **low-dimensional latent space**

CORL2021

Representation: Action Spaces

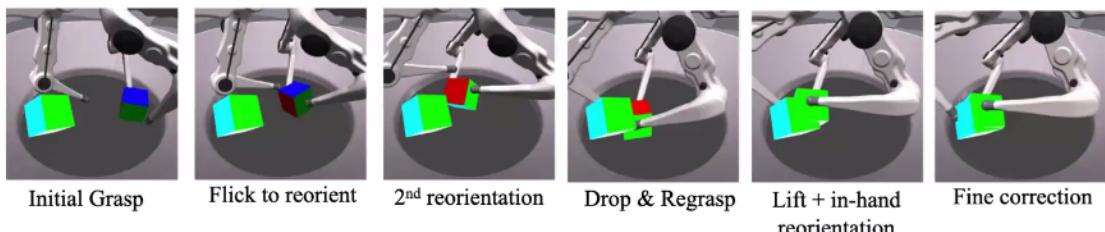
Application: Multi-finger Manip

- Able to train in <24h on 1 GPU



Representation: Action Spaces

Application: Multi-finger Manip

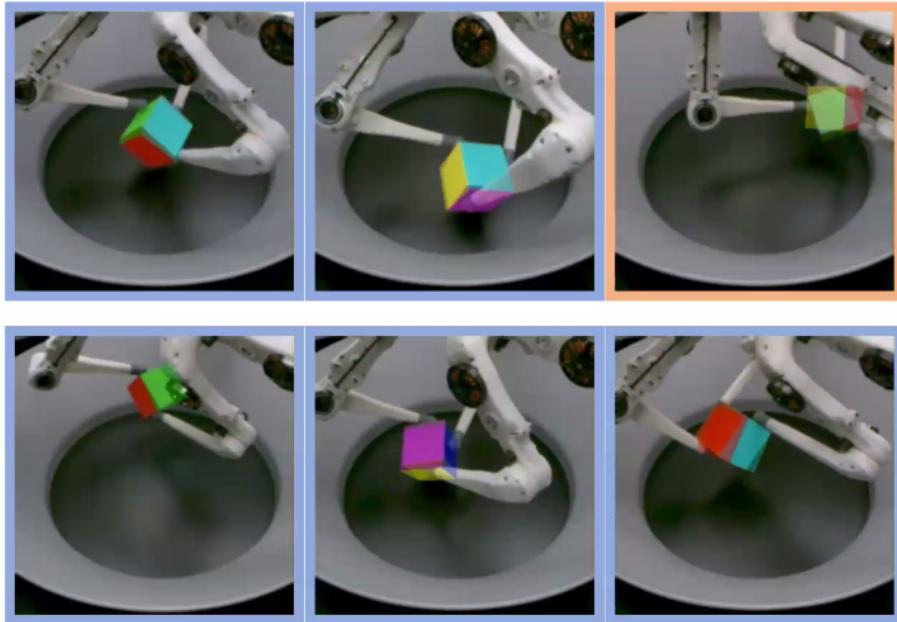


Sim2Real Results

Robotics as a Service

No physical
robot access

83% Success
(Real Time Videos)



Structure

State/Action Reps.
VICES IROS19
LASER ICRA21
Making Sense ICRA19
Unsuo KPs PAMI21

Inductive Biases
C-Learning ICLR21
OCEAN UAI20
D2RL arXiv20

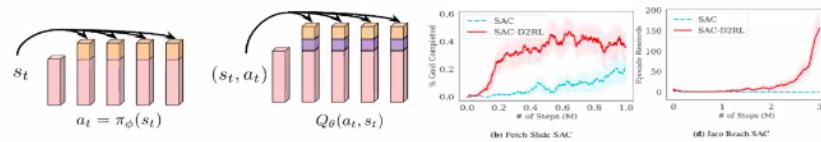
Structure in Planning
CAVIN CoRL20
Skill Hierarchy ICLR21
Finding-IT, CVPR18

Neural Programming
NTP ICRA18
NTG CVPR19
Cont.Relax IROS19

Representations for Robot Learning

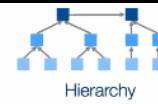
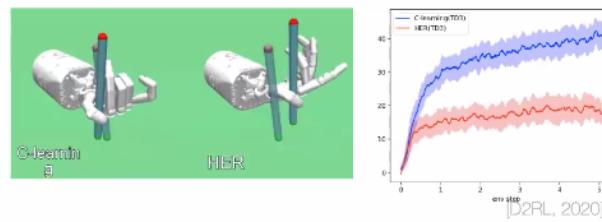
How can better architectures and learning rules lead to generalization in Robotics?

I. Using Dense connections in Policy/Value improves sample efficiency



II. Learning Cumulative Accessibility $C(s, a, h)$ is better than $Q(s, a)$

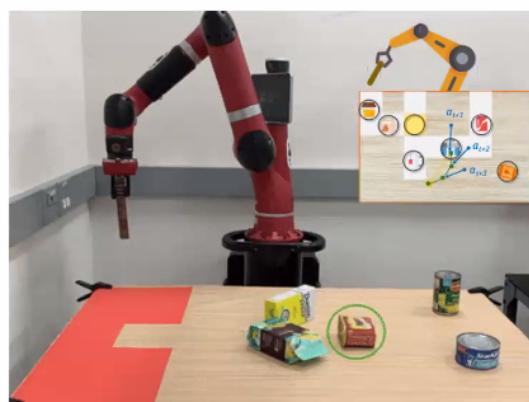
Can represent multimodal horizon aware solutions as well as reachability



Representations for Planning

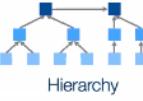
What model structure enables longer term planning?

Hierarchical structure helps in long term discrete and continuous planning



CAVIN, CoRL 2019

Structure



State/Action Reps.
VICES IROS19
LASER ICRA21
Making Sense ICRA19
Unsupo KPs PAM121

Inductive Biases
C-Learning ICLR21
OCEAN UAI20
D2RL arXiv20

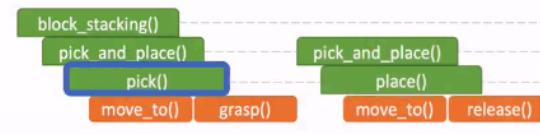
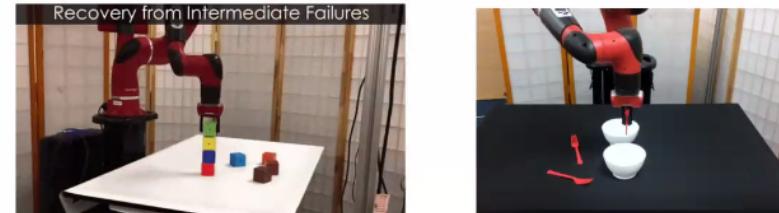
Structure in Planning
CAVIN CORL20
Skill Hierarchy ICLR21
Finding-IT, CVPR18

Neural Programming
NTP ICRA18
NTG CVPR19
Cont.Relax IROS19

Representations for Planning

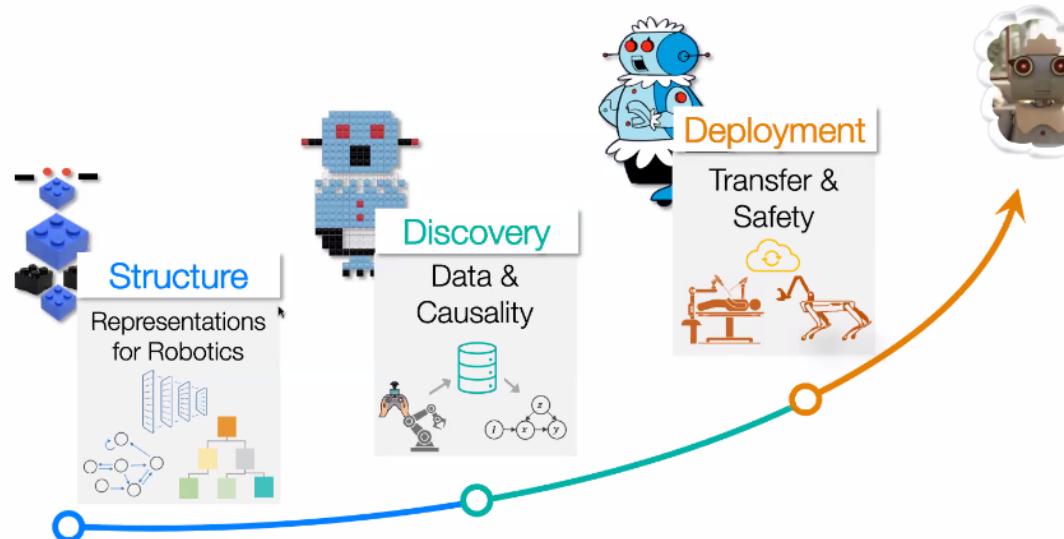
What model structure enables longer term planning?

Program Induction provides a very efficient model of compositional generalization

NTP, ICRA 2018, NTG, CVPR 2019

Building blocks of Generalizable Autonomy in Robotics



Structure
Representations for Robotics

Discovery
Data & Causality

Deployment
Transfer & Safety

Data in Robotics

X Short-Horizon skills
X Skill Specific learning

X Platform dependent data
X Scaling to other skills ?

X Small datasets (minutes)
X Low diversity ?

Manipulation

Mason & Salisbury 1985
Srinivasa et al 2010
Berenson 2013
Odhner1 et al 2014
Chavan-Dafle et al 2014
Yamaguchi, et. al, 2015

Li , Allen et al. 2015
Yahya et al., 2016
Schenck et al. 2017
Mar et al. 2017
Laskey et al 2017
Quispo et al 2018

Grasping

Mishra et al 1987
Ferrari & Cannby, 1992
Ciocarlie & Allen, 2009
Dogar & Srinivasa, 2011
Rodriguez et al. 2012
Bohg et al 2014

Pinto & Gupta, 2016
Levine et al 2016
Mahler et al 2017
Jang et al 2017
Viereck et al 2017
...

Imitation

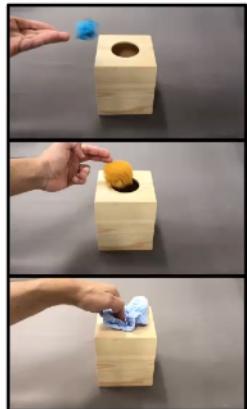
Abbeel et al, 2004
Ratliff et al 2006,
Ziebart et al, 2009
Argall et al, 2009,
Boularias et al., 2011
Montfort et al 2015,
Wulfmeier et al 2015,

Krishnan et al 2017
Finn et al. 2017
Vecerik et al. 2017
Rajeswaran et al 2018
Zhu et al 2018
Ravichandar et al 2020...

Data for Robotics

Embodies Human Cognition and Dexterity

Diversity
Many Solutions



Dexterity
Sophisticated Manipulation



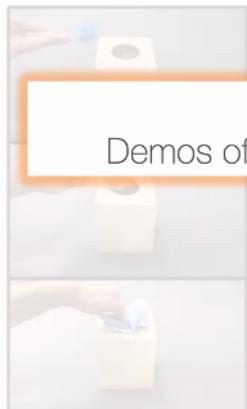
Large-Scale
Many Problem Instances



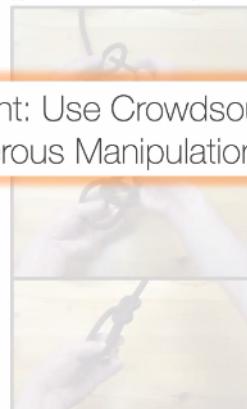
Data for Robotics

Embodies Human Cognition and Dexterity

Diversity
Many Solutions



Dexterity
Sophisticated Manipulation



Large-Scale
Many Problem Instances



Insight: Use Crowdsourcing

Demos of Dexterous Manipulation (not just Labels)



Operator specifies full 6-DoF motion of the arm by moving their phone.

Roboturk Scalability

Multiple Simultaneous Teleoperation Connections



CoRL 2018, IROS 2019

Teleoperation across the world with RoboTurk



RoboTurk in the Swiss Alps
at ~7000 ft



Real-Time Teleoperation from Macau to
Stanford – over 11,000 km!

CoRL 2018, IROS 2019

RoboTurk Pilot Datasets

Simulated Data

137.5 hours of demonstrations

22 hours of total platform usage

3 dexterous manipulation tasks

3224 total attempted demos

15 novice, remote users

Real Robot Data

111 hours of robot demos

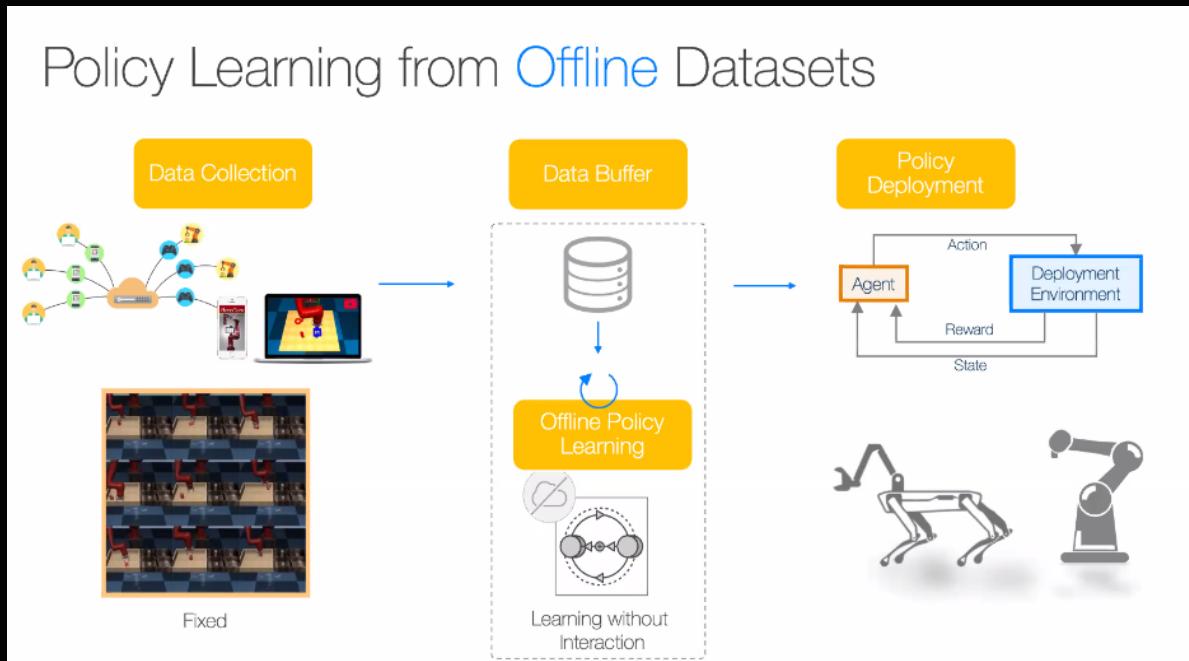
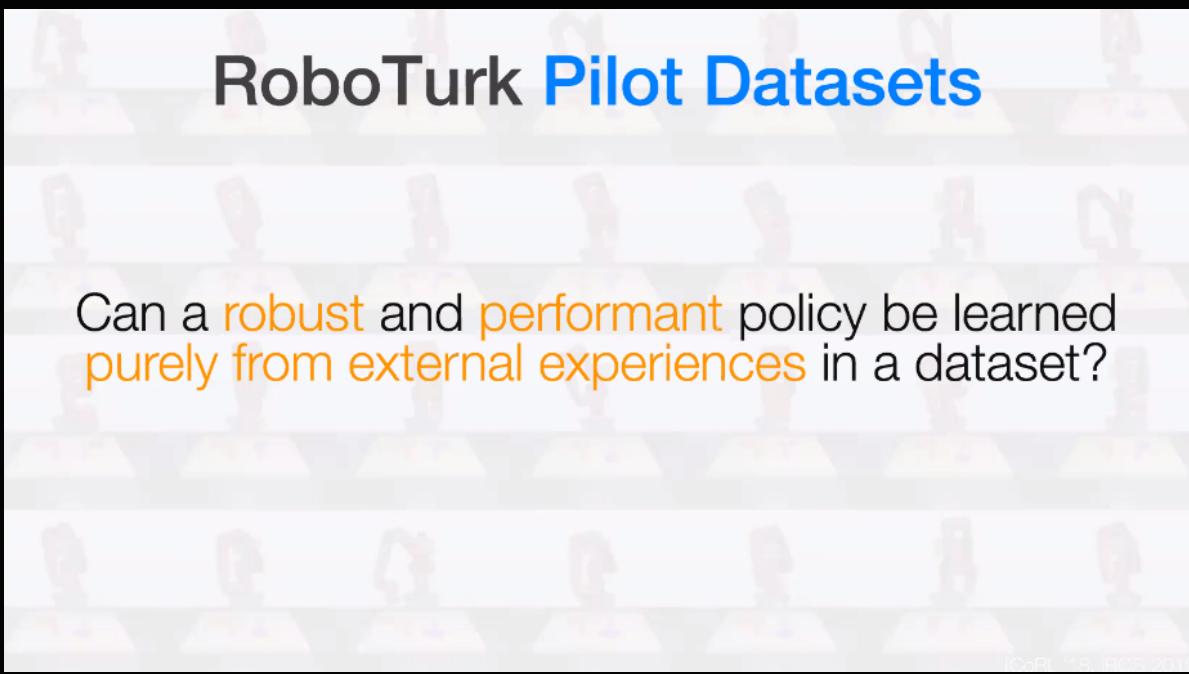
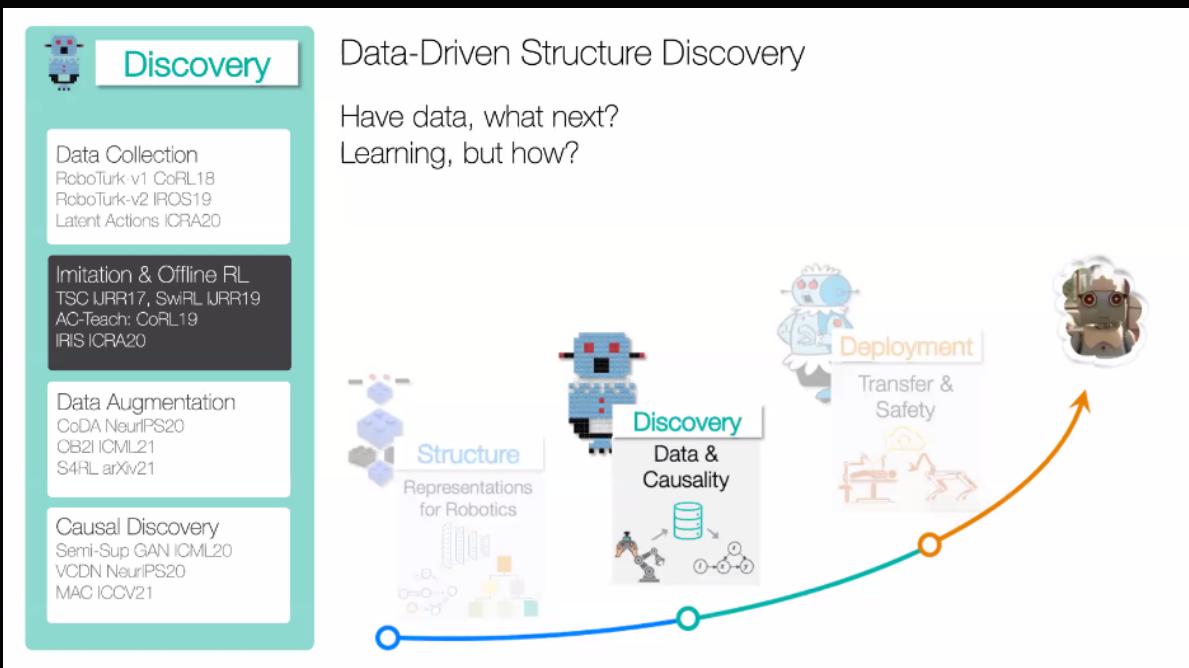
1 week of data collection

3 dexterous manipulation tasks

2144 total demonstrations

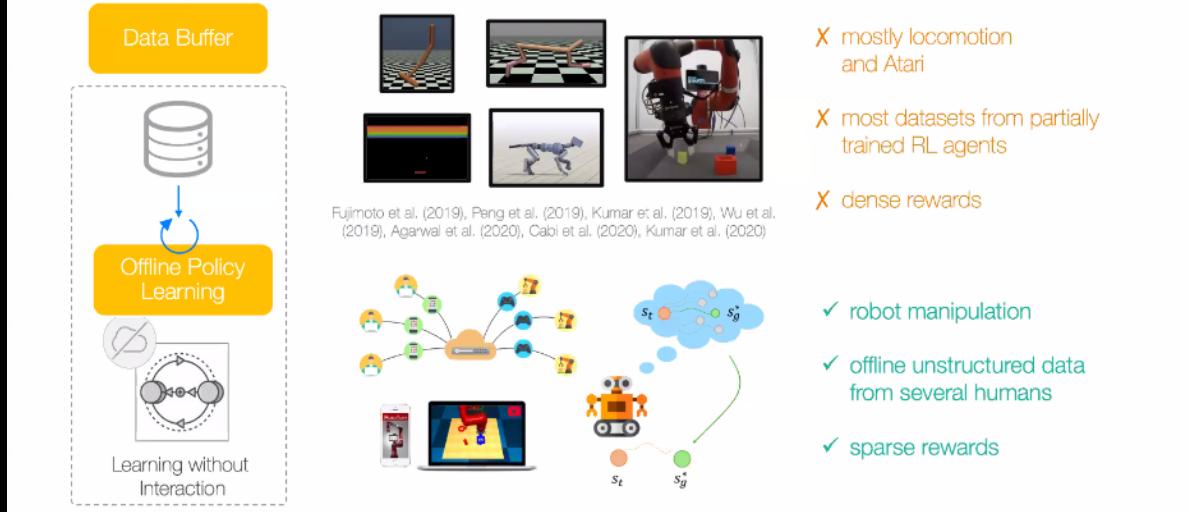
54 non-expert users

CoRL 2018, IROS 2019



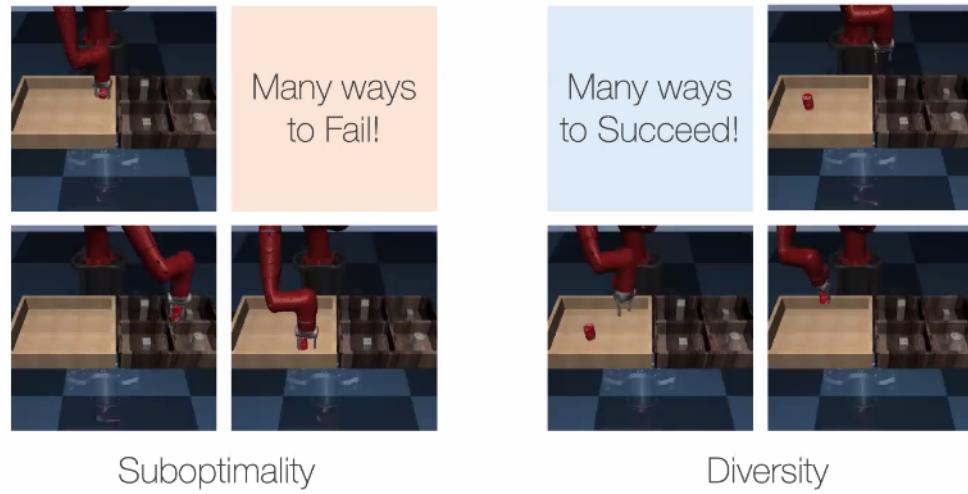
Offline RL

Challenges of Manipulation Datasets



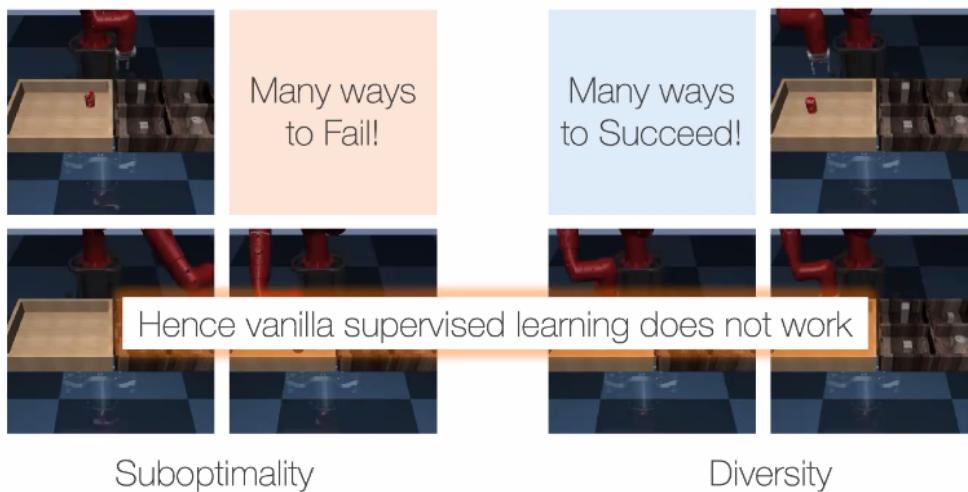
Offline RL

Challenges of Large Datasets



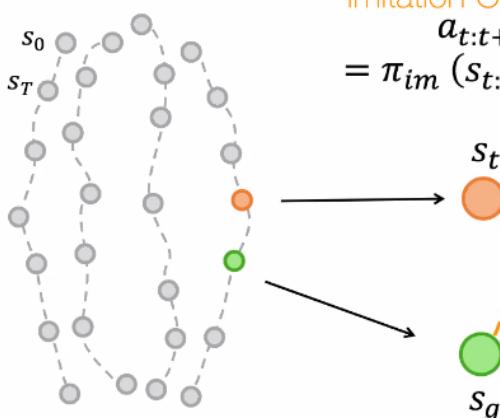
Offline RL

Challenges of Large Datasets



Offline RL: IRIS

Demonstrations



Low-Level
Goal-Conditioned
Imitation Controller

$$a_{t:t+T} = \pi_{im}(s_{t:t+T} | s_g)$$

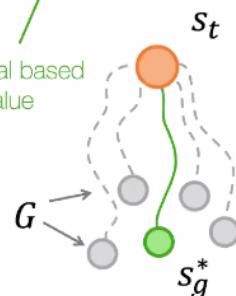
Goal Proposals

Select goal based
on value

High-Level
Goal Selection
Mechanism

$$G = \{s_g^i \sim D_g(s_t)\}_{i=1}^{n_g}$$

$$s_g^* = \max_{s_g \in G} V(s_g)$$



IRIS ICRA 2020

Offline RL: IRIS Evaluation

Behavioral Cloning (BC)



Batch-Constrained Q-Learning (BCQ)



IRIS (ours)



Unable to imitate reasonable behavior from dataset
since models do not account for diverse data.

Goal-conditioned imitation is critical for
learning from diverse dataset.

IRIS ICRA 2020

Offline RL: IRIS Evaluation

Goal Selection Mechanism



Goal-Conditioned Imitation

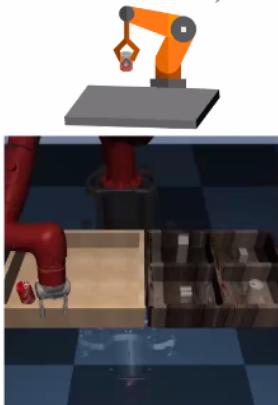


The low-level controller reaches the goals set by the high-level to accomplish the task.

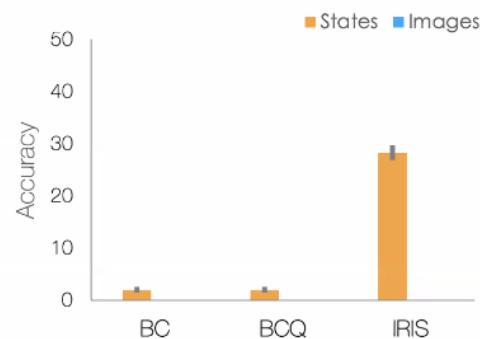
IRIS ICRA 2020

Offline RL: IRIS Evaluation

Ground Truth Robot/Object States



IRIS with States



IRIS ICRA 2020

Offline RL: IRIS Evaluation

Ground Truth Robot/Object States

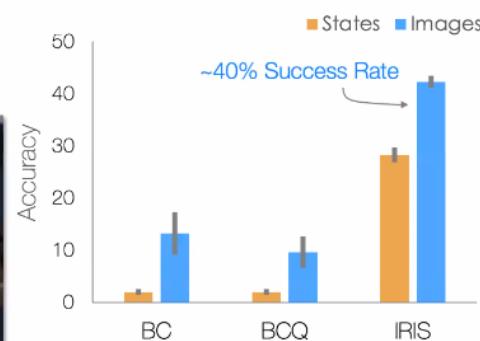


IRIS with States

Learned Keypoint Representations



IRIS with Images



IRIS ICRA 2020



Discovery

Data Collection
RoboTurk v1 CoRL18
RoboTurk-v2 IROS19
Latent Actions ICA20

Imitation & Offline RL
TSC IJRR17, SwiRL IJRR19
AC-Teach: CoRL19
IRIS ICRA20

Data Augmentation
CoDA NeurIPS20
CB2I CML21
S4RL arXiv21

Causal Discovery
Semi-Sup GAN ICML20
VCDN NeurIPS20
MAC ICCV21

Data-Driven Structure Discovery

Collect Structured Dexterous Data
Learn with Offline RL

Is data enough? Can we use it efficiently?

Structure
Representations
for Robotics



Deployment

Transfer &
Safety



Data Augmentation in RL

Do more with the same data



Left Arm Pick and Place



Right Arm Pick and Place

Which of the following is possible (only based on observed data)



Independent
Compositional Generalization

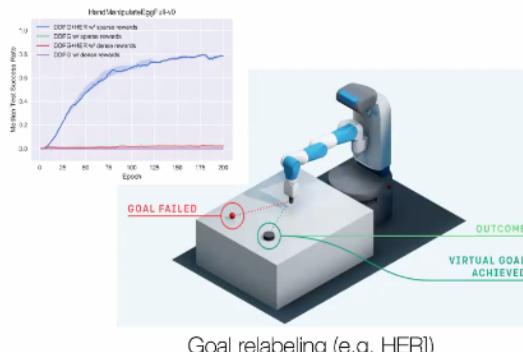


Not-Independent (!)
Hence need evidence of possibility

Data Augmentation in RL

How to do this Algorithmically

- Substantial performance boosts!



Goal relabeling (e.g., HER)]

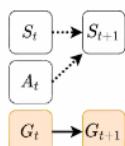
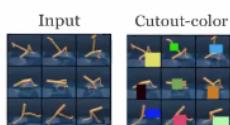


Visual Input relabeling (e.g., RAD)]

RAD: Laskin et al 2020, DrQ: Kostrikov et al 2020, CURL: Laskin et al. 2020, S4RL: Sinha et al. 2021, HER: Andrychowicz et al 2018

Data Augmentation in RL

Unified View

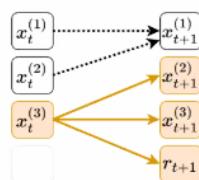


Goal is independent of
State/Action Dynamics

Visual characteristics (e.g.,
crop) are independent of
physical dynamics

Counterfactual reasoning to
generate new, causally valid
(counterfactual) data!

Generic CoDA



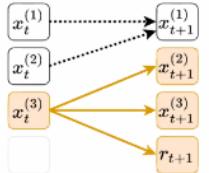
Given two independent
mechanisms, Relabel one
(conditional independence!)

RAD: Laskin et al 2020, DrQ: Kostrikov et al 2020, CURL: Laskin et al. 2020, S4RL: Sinha et al. 2021, HER: Andrychowicz et al 2018

Counterfactual Data Augmentation

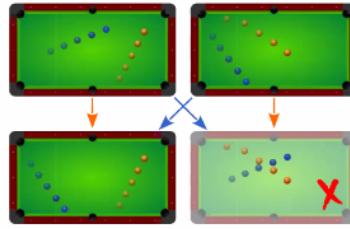
Counterfactual reasoning to generate new, causally valid (counterfactual) data!

Generic CoDA



Given two independent mechanisms, Relabel one (conditional independence!)

- ✓ Model-Free relabelling
- ✗ But Causal Independence is not Global

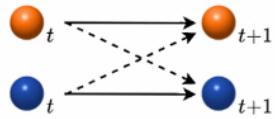


- For the most part, balls behave independently, and we can use CoDA
- But balls are not always independent, so this can also produce nonsense

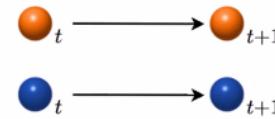
CoDA NeurIPS 2020

Counterfactual Data Augmentation

Local Causal Model



Global Model



Local Model

$$\mathcal{M}_t = \langle V_t, U_t, \mathcal{F} \rangle \xrightarrow{\text{Condition on } (s_t, a_t) \in \mathcal{L}} \mathcal{M}_t^{\mathcal{L}} = \langle V_t^{\mathcal{L}}, U_t^{\mathcal{L}}, \mathcal{F}^{\mathcal{L}} \rangle$$

Structural Causal Model (SCM) that marginalizes across all possible transitions

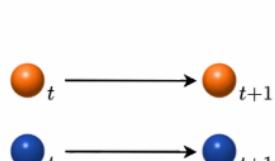
Local Causal Model (LCM) that behaves like the global SCM in local subspace \mathcal{L}

CoDA NeurIPS 2020

Counterfactual Data Augmentation

Learning Local Causal Model

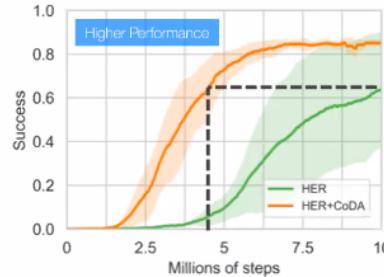
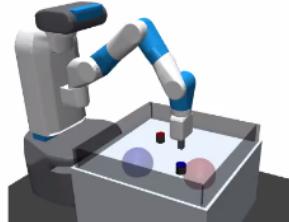
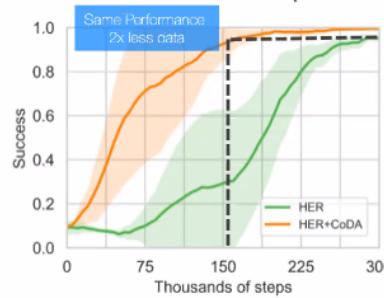
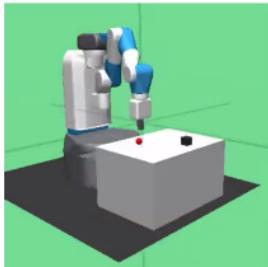
- Input: 2 balls, each with 4 features: $[x, y, \dot{x}, \dot{y}]$
 - [[1.23, -0.73, 1.31, 1.07], [-0.6, 2.51, -1.51, -0.89]]
- Output: Adjacency matrix M of the causal graph (between x_t and x_{t+1})
- (intuition) M: the input-output Jacobian is non-zero



x_{t+1} x_{t+1}
 x_t x_t
 $[[1., 1., 1., 1., 0., 0., 0., 0.],$
 $[1., 1., 1., 1., 0., 0., 0., 0.],$
 $[1., 1., 1., 1., 0., 0., 0., 0.],$
 $[1., 1., 1., 1., 0., 0., 0., 0.],$
 $[0., 0., 0., 0., 1., 1., 1., 1.],$
 $[0., 0., 0., 0., 1., 1., 1., 1.],$
 $[0., 0., 0., 0., 1., 1., 1., 1.],$
 $[0., 0., 0., 0., 1., 1., 1., 1.]]$

CoDA NeurIPS 2020

CoDA: Goal-Conditioned (Online) RL



CoDA NeurIPS 2020



Discovery

Data Collection
RoboTurk-v1 CoRL18
RoboTurk-v2 IROS19
Latent Actions ICRA20

Imitation & Offline RL
TSC URR17, SwRL IJRR19
AC-Teach: CoRL19
IRIS ICRA20

Data Augmentation
CoDA NeurIPS20
CB2I ICML21
S4RL arXiv21

Causal Discovery
Semi-Sup GAN ICML20
VCDN NeurIPS20
MAC ICCV21

Data for Robotics

How to collect large-scale structured supervision and then learn from it?

Can we use the data better?

How to learn this structure automatically?



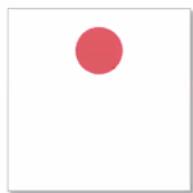
Deployment

Transfer & Safety

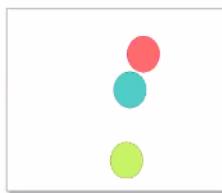


Causal Discovery

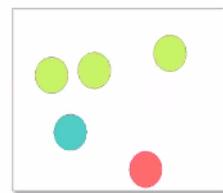
Guess Model



1-Ball bouncing on Flat surface with no damping



3-Balls (different) in a box with no damping



5-Balls (different) in a box with no damping



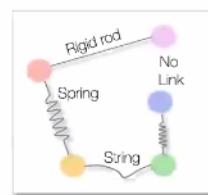
5-Balls (different) in a box with unknown connections

Model is **tricky** to Guess

- Latent Mechanisms impose constraints

Can't **Memorize**, need to **Generalize** to new latent graphs

- Each Test case may be a different graph

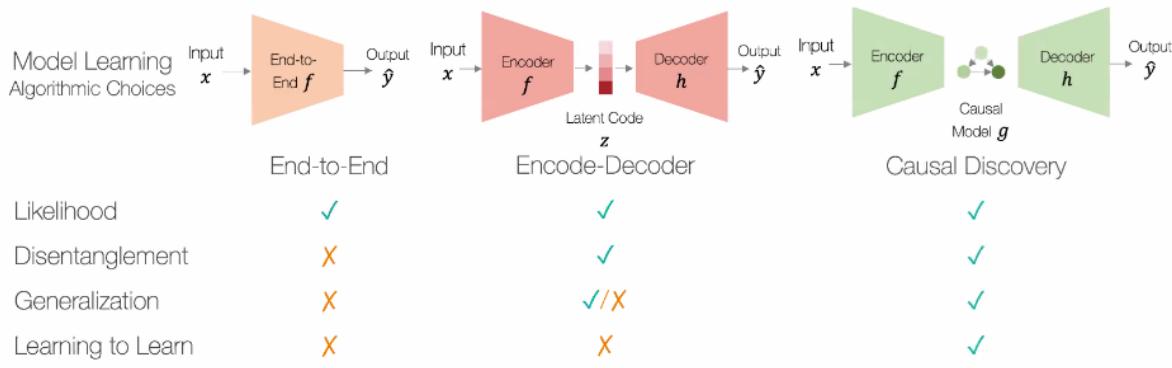


Causal Discovery



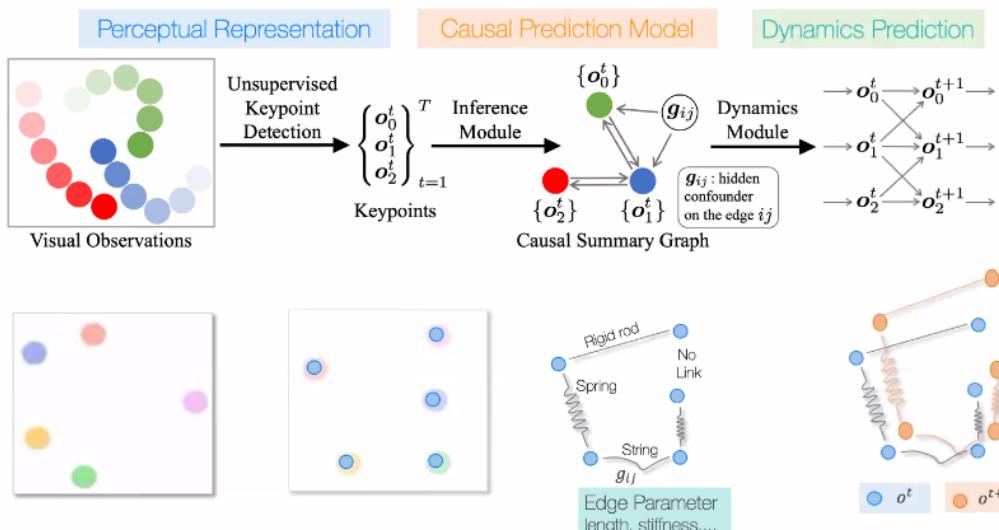
Causal Discovery of Dynamics Learning Latent Generative Dynamics Model

Output: Single model to multiple latent generative graphs



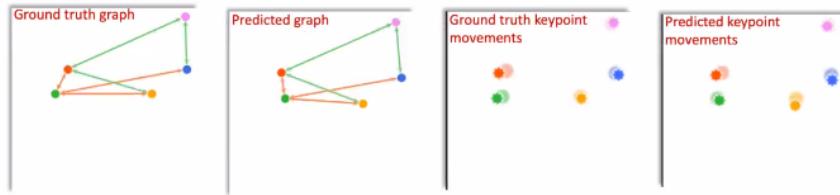
VCDN NeurIPS 2020

Visual Causal Discovery Network



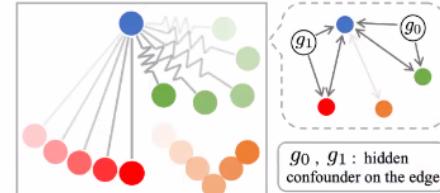
VCDN NeurIPS 2020

Visual Causal Discovery Network Evaluation



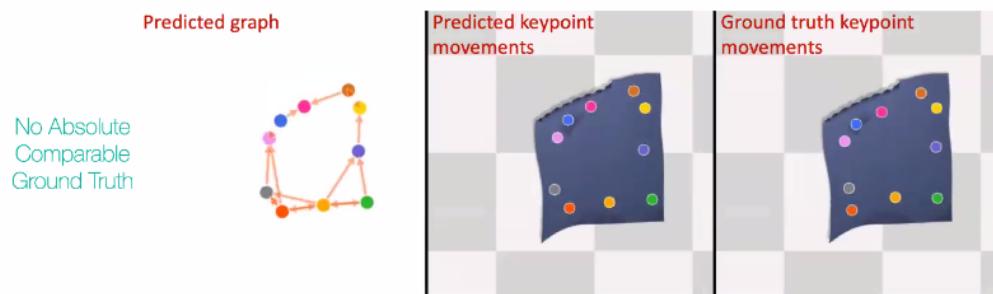
Balls Environment

- N-balls with different edge types {spring, rigid, null}
- Each edge has random parameters
- Dataset: 4800 Training Episodes with random latent generative model



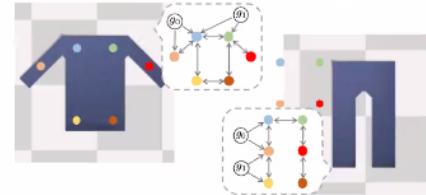
VCDN NeurIPS 2020

Visual Causal Discovery Network Evaluation



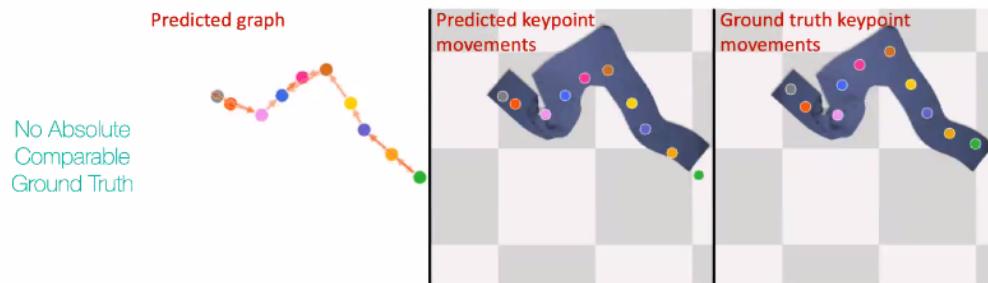
Fabrics Environment

- 3 Types: Shirts, Towels, Pants
- Each instance has variable geometry
- Dataset: 1800 Training Episodes with random latent generative model



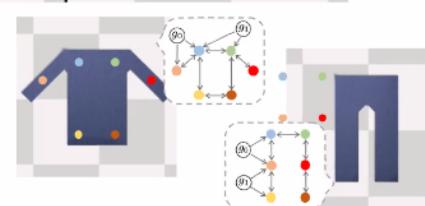
VCDN NeurIPS 2020

Visual Causal Discovery Network Evaluation



Fabrics Environment

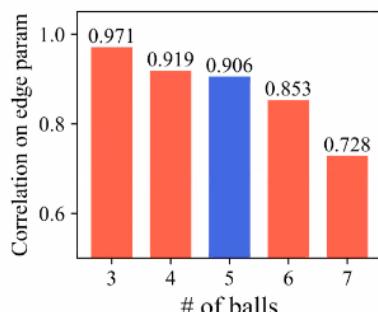
- 3 Types: Shirts, Towels, Pants
- Each instance has variable geometry
- Dataset: 1800 Training Episodes with random latent generative model



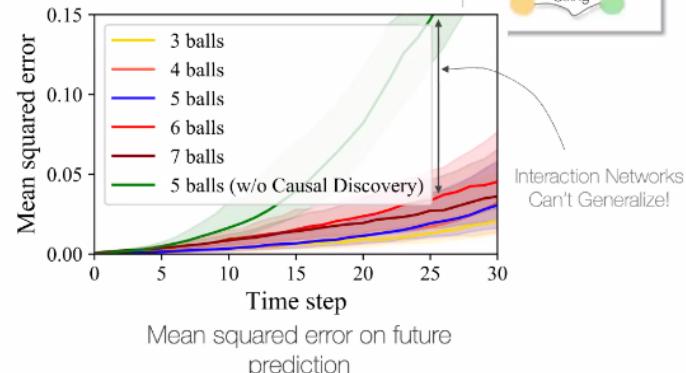
VCDN NeurIPS 2020

Visual Causal Discovery Network

Extrapolation Generalization



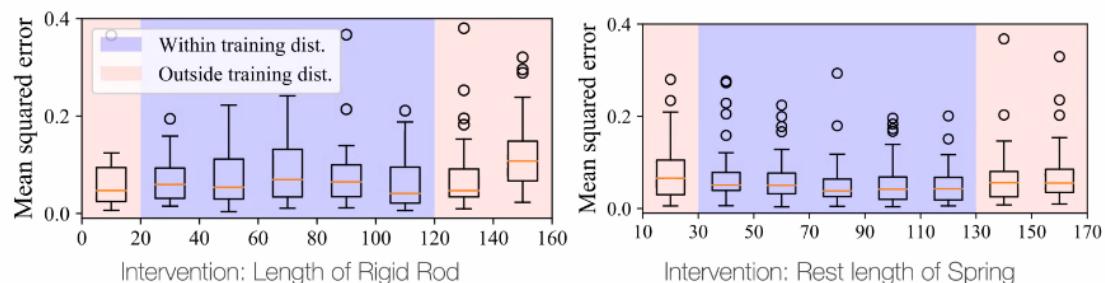
Correlation on the rest length of the spring relation
(No Ground Truth for training)



VCDN NeurIPS 2020

Visual Causal Discovery Network

Counterfactual Generalization

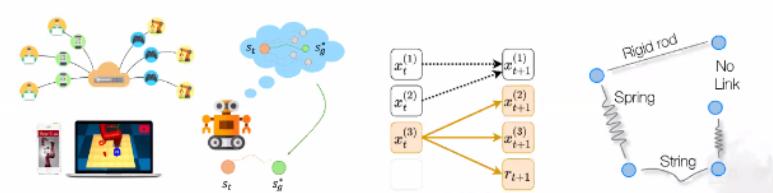


Change specific underlying parameters for given datapoints

VCDN NeurIPS 2020



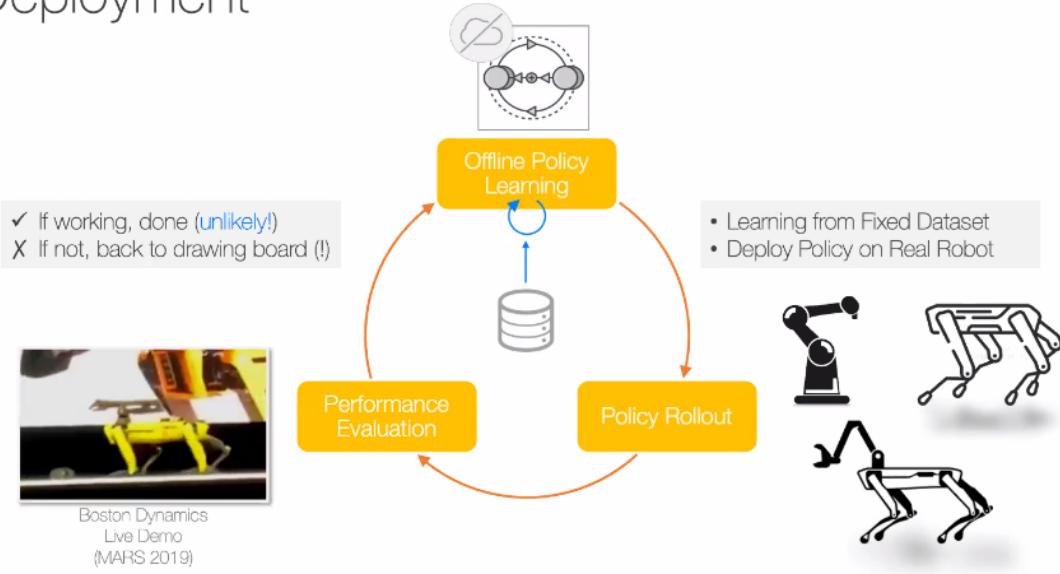
Data-Driven Structure Discovery



Structured-supervision through scalable crowdsourcing
Causal Discovery is vital for compositional generalization



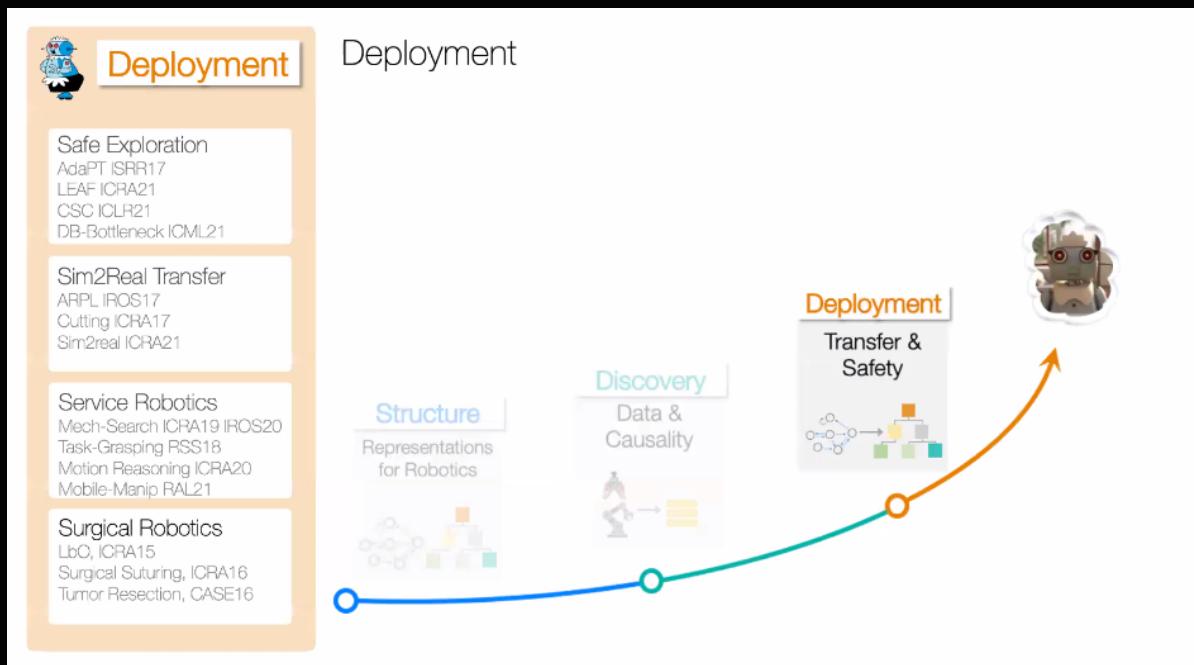
Deployment



Deployment: Mission Critical Systems



Need to evaluate behavior before deployment and safety during deployment!



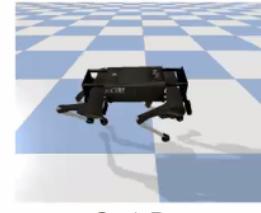
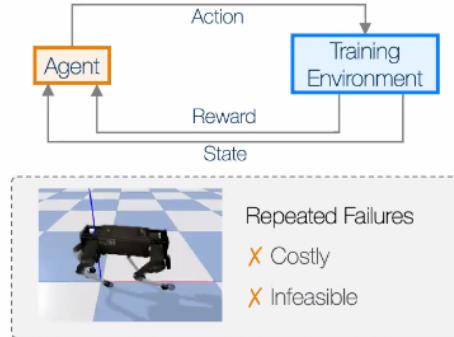
Deployment

Challenge: Safe Exploration

How should an agent learn without **Catastrophic Failures** during training?



Initial State



Goal: Run

Conservative Safety Critics ICLR 2021

Deployment

Challenge: Safe Exploration

How should an agent learn without **Catastrophic Failures** during training?

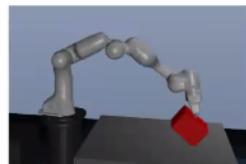
Respecting constraints of the form $\mathbb{P}(\text{failure}) \leq \chi$



Dexterous Manipulation



Objective: Learn to Push
Failure: Glass Topples



Legged Locomotion



Objective: Learn to Run
Failure: Fall Down



Conservative Safety Critics ICLR 2021

Deployment

Challenge: Safe Exploration

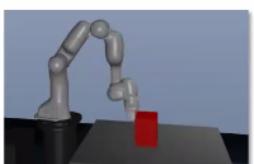
How should an agent learn without **Catastrophic Failures** during training?

Respecting constraints of the form $\mathbb{P}(\text{failure}) \leq \chi$

- Assume a binary constraint function $\mathbb{I}(\text{fail})$ from Environment

$$C: \mathcal{S} \rightarrow \{0, 1\}, (C \equiv \mathbb{I}\{\text{failure}\})$$

- Analogous to sparse rewards {+1: failure occurs, 0: otherwise}



Dexterous Manipulation



Objective: Learn to Push
Failure: Glass Topples

$$C = 1 \text{ if } \mathbb{I}(\text{topple})$$



Legged Locomotion



Objective: Learn to Run
Failure: Fall Down

$$C = 1 \text{ if } \mathbb{I}(\text{fall})$$

Conservative Safety Critics ICLR 2021

Conservative Safety Critic

Insight: Learning Constraints as Critics

Constrained MDP Objective:

$$\begin{aligned} \pi^* &= \underset{\pi \in \Pi_C}{\operatorname{argmax}} V_R^\pi(\mu) && \text{Task Objective} \\ \text{s.t. } \mathbb{P}(\text{failure} | \mu) &\leq \chi && \text{Safety Constraint} \end{aligned}$$

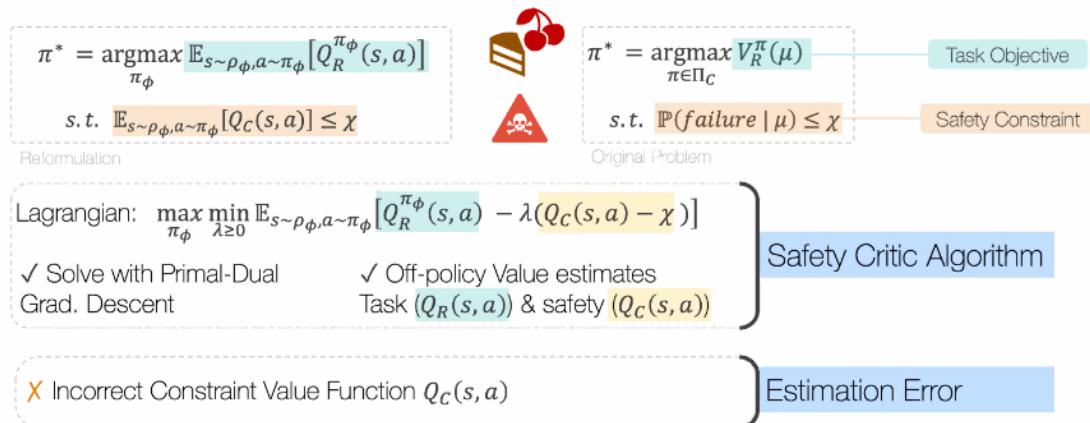
$$\begin{aligned} \mathbb{P}(\text{failure} | \mu) &= \mathbb{E}_{\tau \sim \pi} [\mathbb{I}\{\text{failure}\}] \\ &= \mathbb{E}_{\tau \sim \pi} \left[\sum_{t=0}^{\infty} C(s_t) \right] && \text{Expectation over observed Rollouts} \\ &= V_C^\pi(\mu) && \text{Safety Value Function} \end{aligned}$$

Conservative Safety Critics ICLR 2021

Conservative Safety Critic

Insight: Learning Constraints as Critics

Constrained MDP Objective:



Conservative Safety Critics ICLR 2021

Conservative Safety Critic

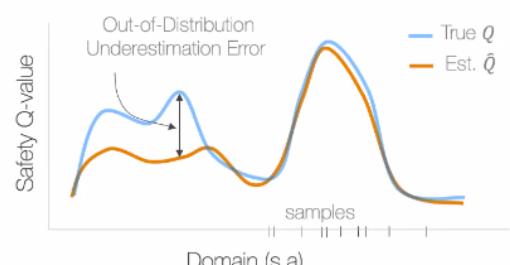
Insight-2: Be more Cautious with Overestimation

✗ How to fix Estimation error in $Q_C(s, a)$

Q-Learning (review)

$$Q_C(s, a) \leftarrow r(s, a) + \mathbb{E}_{a' \sim \pi_{new}} [Q_C(s', a')] - Q_{Target}$$

Objective: $\min_Q \mathbb{E}_{a \sim \pi_B} [Q_C(s', a') - Q_{Target}]$



Objective: $\min_Q \mathbb{E}_{(s,a) \sim Data} [Q_C(s', a') - Q_{Target}] - \alpha \mathbb{E}_{s \sim Data, a \sim \pi_\phi} [Q_C(s', a')]$

Regularizer (Push Q-values up)

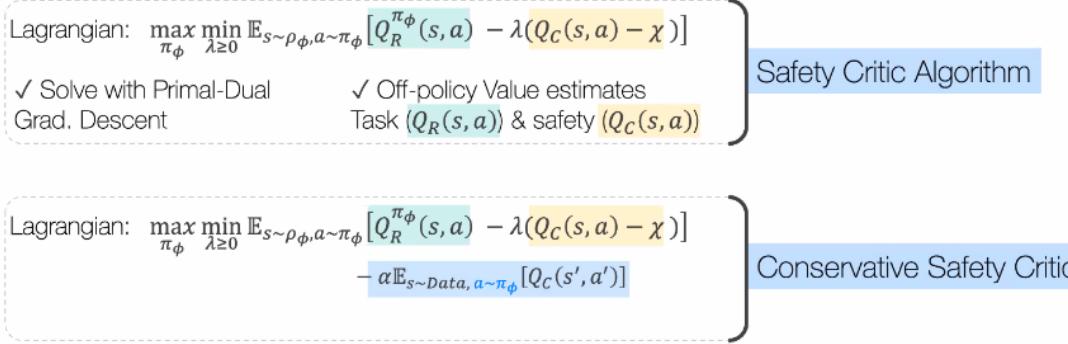
Conservative Q-Learning

(CQL, Kumar et al 2020)

Provably $\hat{Q} \geq Q$

Conservative Safety Critics ICLR 2021

Conservative Safety Critic

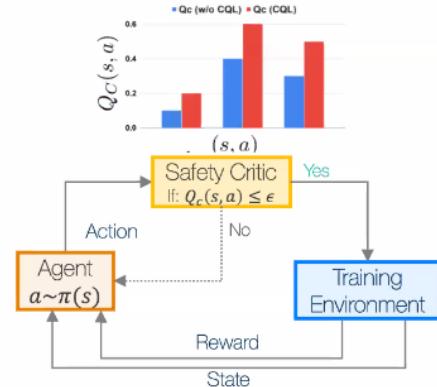


Conservative Safety Critics ICLR 2021

Conservative Safety Critic Inference Mechanism



Vanilla RL



Conservative Safety Critic

Lagrangian: $\max_{\pi_\phi} \min_{\lambda \geq 0} \mathbb{E}_{s \sim \rho_\phi, a \sim \pi_\phi} [Q_R^{\pi_\phi}(s, a) - \lambda(Q_C(s, a) - \chi)] - \alpha \mathbb{E}_{s \sim \text{Data}, a \sim \pi_\phi} [Q_C(s', a')]$

Conservative Safety Critics ICLR 2021

Conservative Safety Critic Guarantees

Safe Training: Constraint is satisfied during for each iterate

Theorem 1. Consider policy updates that solve the constrained optimization problem defined in Equation 5. With high probability $\geq 1 - \omega$, we have the following upper bound on expected probability of failure $V_C^{\pi_{\phi_{new}}}(\mu)$ for $\pi_{\phi_{new}}$ during every policy update iteration:

$$V_C^{\pi_{\phi_{new}}}(\mu) \leq \chi + \zeta - \frac{\Delta}{1-\gamma} + \frac{\sqrt{2\delta}\gamma\epsilon_C}{(1-\gamma)^2} \quad \text{where} \quad \zeta \leq \frac{C'\sqrt{\log(1/\omega)}}{|N|} \quad (7)$$

Safety not just at the end but all along training

Sample Efficiency: Bounded effect on rate of convergence

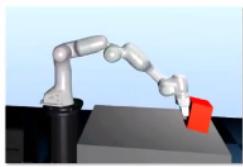
Theorem 2 (Convergence rate for policy gradient updates with the safety constraint). If we run the policy gradient updates through equation 5, for policy π_ϕ , with μ as the starting state distribution, with $\phi^{(0)} = 0$, learning rate $\eta > 0$, and choose α as mentioned in the discussion of Theorem 1, then for all policy update iterations $T > 0$ we have, with probability $\geq 1 - \omega$,

$$V_R^*(\mu) - V_R^{(T)}(\mu) \leq \frac{\log |\mathcal{A}|}{\eta T} + \frac{1}{(1-\gamma)^2 T} + K \frac{\sum_{t=0}^{T-1} \lambda^{(t)}}{\eta T} \quad \text{where} \quad K \leq (1-\chi) + \frac{4\sqrt{2\delta}\gamma}{(1-\gamma)^2}$$

Adding constraint does not affect sample efficiency

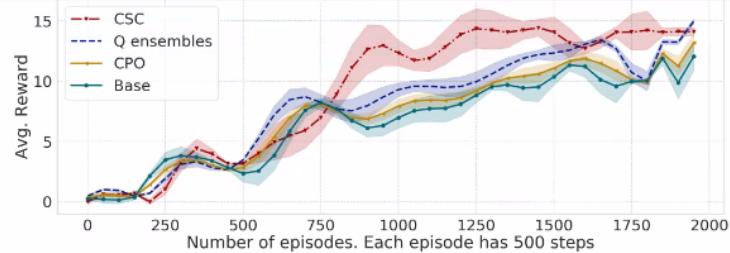
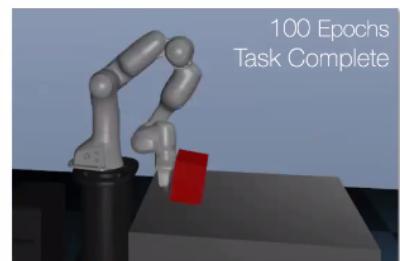
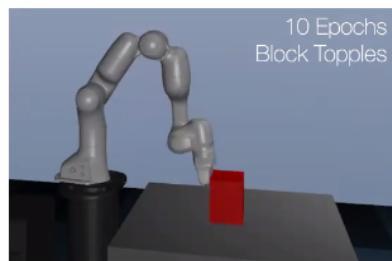
Conservative Safety Critics ICLR 2021

Conservative Safety Critic Evaluation



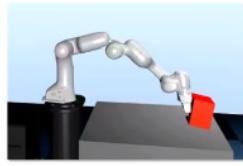
Dexterous Manipulation

Objective: Learn to Push
Failure: Glass Topple



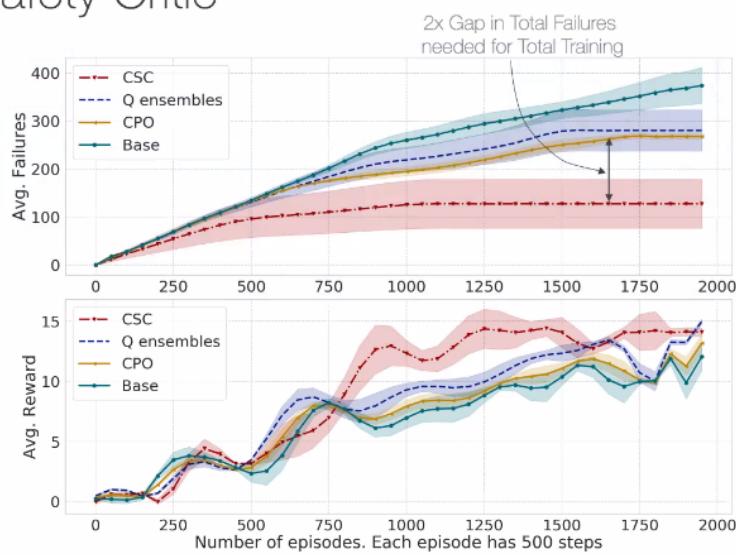
Conservative Safety Critics ICLR 2021

Conservative Safety Critic Evaluation



Dexterous Manipulation

Objective: Learn to Push
Failure: Glass Topple



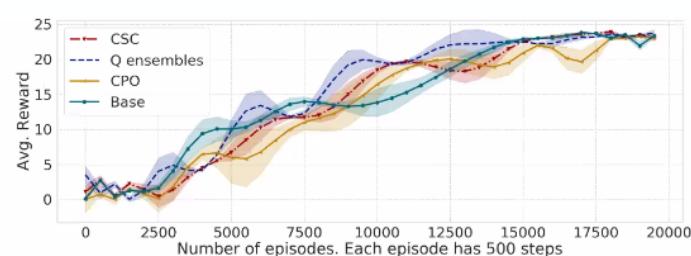
Conservative Safety Critics ICLR 2021

Conservative Safety Critic Evaluation



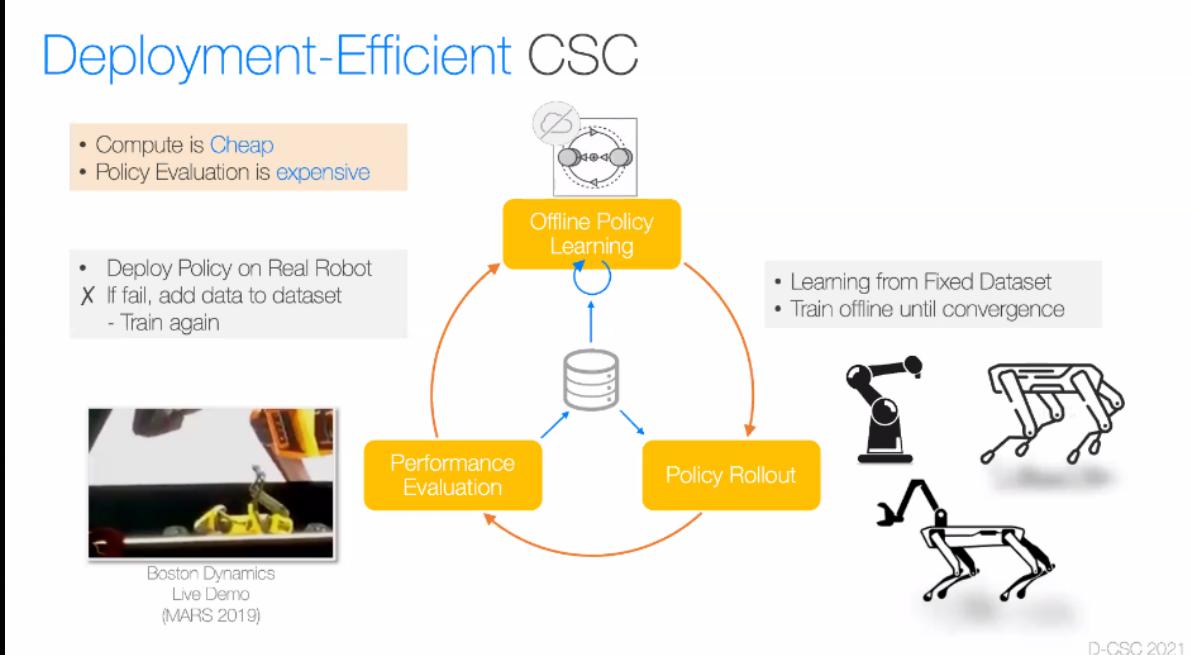
Legged Locomotion

Objective: Learn to Run
Failure: Fall Down

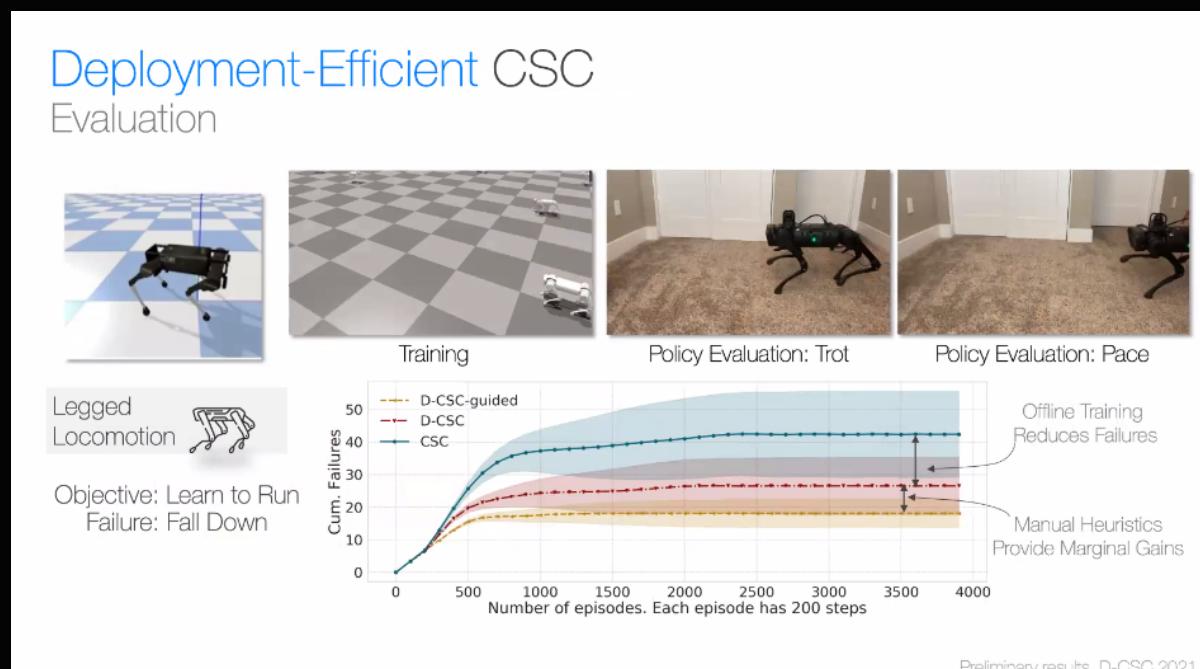


Conservative Safety Critics ICLR 2021

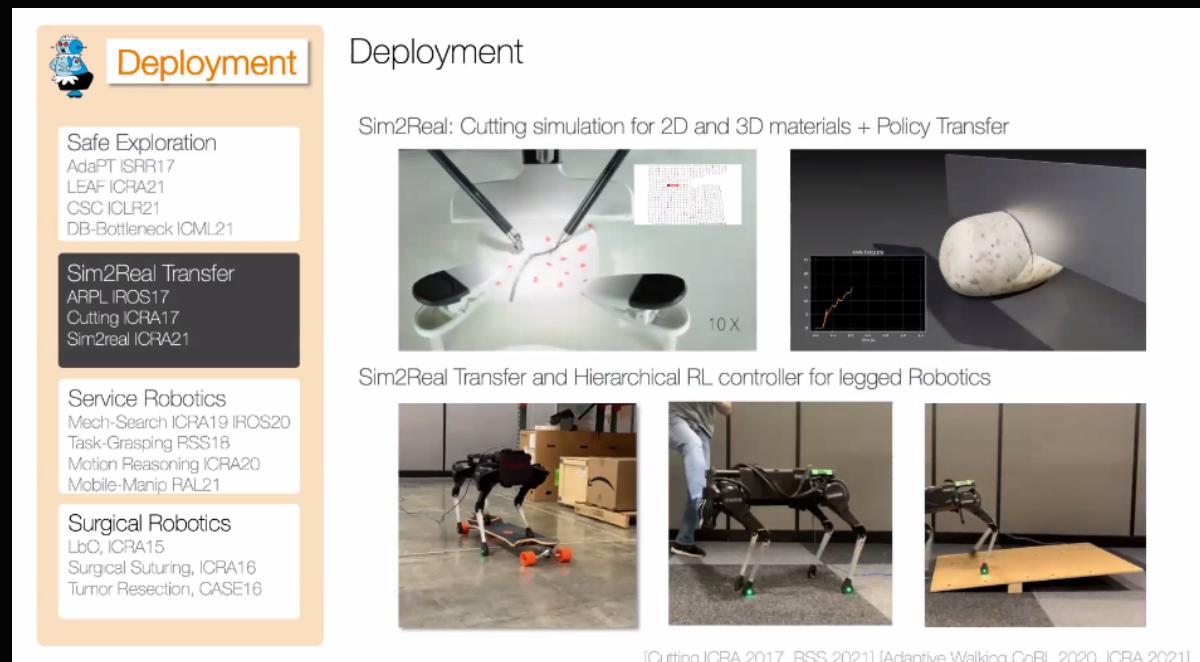
Deployment-Efficient CSC



Deployment-Efficient CSC Evaluation



Deployment





Deployment

Safe Exploration
AdaPT ISRR17
LEAF ICRA21
CSC ICLR21
DB-Bottleneck ICML21

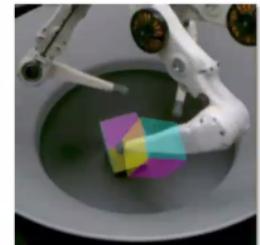
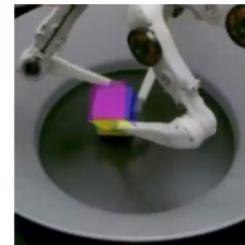
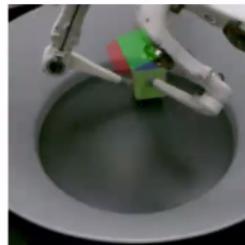
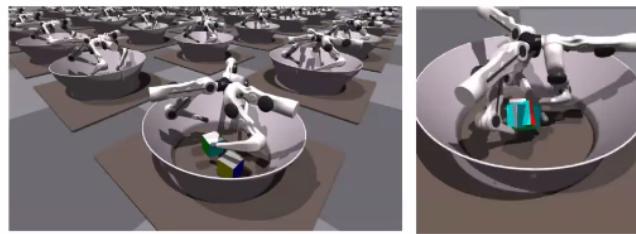
Sim2Real Transfer
ARPL IROS17
Cutting ICRA17
Sim2real ICRA21

Service Robotics
Mech-Search ICRA19 IROS20
Task-Grasping RSS18
Motion Reasoning ICRA20
Mobile-Manip RAL21

Surgical Robotics
LbO, ICRA15
Surgical Suturing, ICRA16
Tumor Resection, CASE16

Deployment

Sim2Real: In-hand manipulation



[CORL 2021 (under review)]



Deployment

Safe Exploration
AdaPT ISRR17
LEAF ICRA21
CSC ICLR21
DB-Bottleneck ICML21

Sim2Real Transfer
ARPL IROS17
Cutting ICRA17
Sim2real ICRA21

Service Robotics
Mech-Search ICRA19 IROS20
Task-Grasping RSS18
Motion Reasoning ICRA20
Mobile-Manip RAL21

Surgical Robotics
LbO, ICRA15
Surgical Suturing, ICRA16
Tumor Resection, CASE16

Deployment

Learning Manipulation Skills

Object Sorting



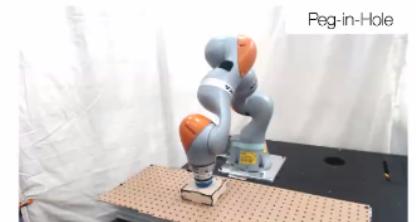
Object Search



Goal Conditioned Imitation



Robot Execution



[Sorting ICRA 2018], [Object Search: ICRA 2019], [Goal Imitation ICRA 2020], [Tool Use RSS2018, RSS 2021], [Peg-in-Hole ICRA 2019]



Deployment

Safe Exploration
AdaPT ISRR17
LEAF ICRA21
CSC ICLR21
DB-Bottleneck ICML21

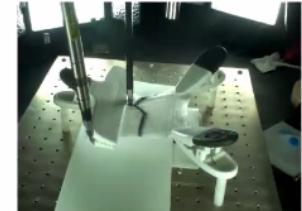
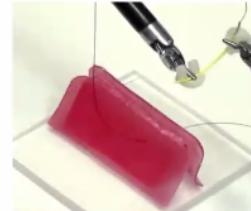
Sim2Real Transfer
ARPL IROS17
Cutting ICRA17
Sim2real ICRA21

Service Robotics
Mech-Search ICRA19 IROS20
Task-Grasping RSS18
Motion Reasoning ICRA20
Mobile-Manip RAL21

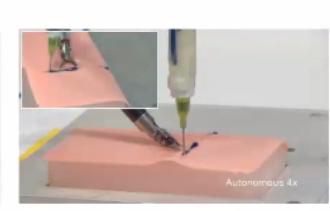
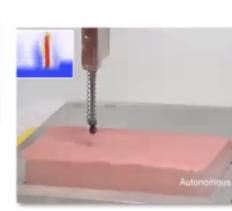
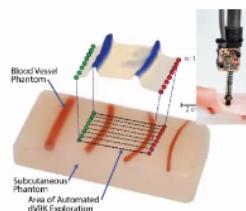
Surgical Robotics
LbO, ICRA15
Surgical Suturing, ICRA16
Tumor Resection, CASE16

Deployment

Surgical Subtask Automation: Imitation Learning for Structured Tasks



Tumor Localization and Resection: Bayesian Search with Noisy Sensors



[Subtask Automation ICRA 2015, ICRA 2016, IJRR 2019] [Tumor Localization CASE2016, Hamlyn 2015]



Deployment

Safe Exploration

AdaPT ISRR17
LEAF ICRA21
CSC ICLR21
DB-Bottleneck ICML21

Sim2Real Transfer

ARPL IROS17
Cutting ICRA17
Sim2real ICRA21

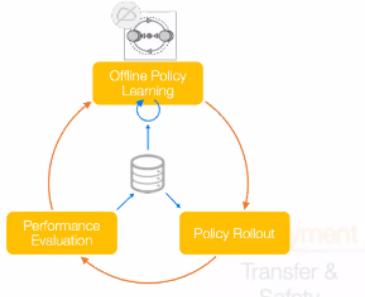
Service Robotics

Mech-Search ICRA19 IROS20
Task-Grasping RSS18
Motion Reasoning ICRA20
Mobile-Manip RAL21

Surgical Robotics

LbO, ICRA15
Surgical Suturing, ICRA16
Tumor Resection, CASE16

Deployment



Practical challenges in Robot Deployment
Novel Algorithms + Systems' view of Robotics



Acknowledgements



Collaborators and Colleagues



Structure

State/Action Reps.
VICES IROS19
LASER ICRA21
Making Sense ICRA19
Unsup KPs PAMI21

Inductive Biases
C-Learning ICLR21
OCCEAN UAI20
D2RL arXiv20

Structure in Planning
CAVIN CORL20
Skill Hierarchy ICLR21
Finding-IT, CVPR18

Neural Programming
NTP ICRA18
NTG CVPR19
Cont.Relax IROS19



Discovery

Data Collection
RoboTurk-v1 CoRL18
RoboTurk-v2 IROS19
Latent Actions ICRA20

Imitation & Offline RL
TSC IJRR17, SwrL IJRR19
AC-Teach: CoRL19
IRIS ICRA20

Data Augmentation
CoDA NeurIPS20
OB2I ICML21
S4RL arXiv21

Causal Discovery
Semi-Sup GAN ICML20
VCDN NeurIPS20
MAC ICCV21



Deployment

Safe Exploration
AdaPT ISRR17
LEAF ICRA21
CSC ICLR21
DB-Bottleneck ICML21

Sim2Real Transfer
ARPL IROS17
Cutting ICRA17
Sim2real ICRA21

Service Robotics
Mech-Search ICRA19 IROS20
Task-Grasping RSS18
Motion Reasoning ICRA20
Mobile-Manip RAL21

Surgical Robotics
LbO, ICRA15
Surgical Suturing, ICRA16
Tumor Resection, CASE16



Agenda

Perception

Dense Action Understanding
Concept Grounding
Learning from Interaction

Planning & Control

Bayes Adaptive RL
Causality via Intervention
Neuro-Symbolic Planning

Data & Benchmarks

Multi-modal Datasets
Taxonomy of Manipulation
Platoon Autonomy

Applications

Personal
Industrial
Healthcare